

The cycle type of two rows of a random Latin square:
exact identities and a reduction of the weak
Cavenagh–Greenhill–Wanless conjecture

[author]

4 October 2026 (second draft of the day) — draft

Abstract

Let L be a uniformly random Latin square of order n and σ the permutation of columns carrying its first row to its second. Cavenagh, Greenhill and Wanless conjectured that the cycle type of σ is asymptotically that of a uniform derangement, in total variation. We prove three exact identities and a reduction theorem: the marked-pair identity, which expresses the ratio of completion counts of adjacent cycle types as the odds of one bit (“rows 1, 2 lie in a common column cycle of a marked column pair”) in a single uniform square; the reduction theorem, by which the conjecture follows from that bit being fair to $o(1/\log n)$; and the ladder identity, by which the bit is exactly fair except on the squares in which no “ladder pair” of rows is flippable. An orbit method bounds the probability of the exceptional set by a product of conditionally independent factors. The conjecture thereby follows from any one of several sufficient conditions, no two of which are known to be comparable except where stated, all concerning the probability of a configuration at rows or columns selected by the square inside the space conditioned on the type of rows 1, 2, and all needed only for types with $O(\log n)$ parts. The reference condition is a “witness lemma” (Hypothesis 5.1): with probability $1 - o(1/\log n)$ some ladder pair of a long ladder has a short row cycle through one marked column avoiding the other. The simplest comes from a third exact identity at the first ladder pair alone, which expresses the bias $\mathbb{P}_X[B] - \mathbb{P}_X[A]$ as an expectation over the event that none of the column pairs along the row permutation of that pair admits a legal toggling turn. That event splits into a short-arc term and a hazard term; we prove that the short-arc term is $o(1/\log n)$, by showing that the per-cell bound of Allsop and Morris for partial Latin squares extends to partial Latin squares containing two complete rows, so that the conjecture follows from the hazard term alone: with probability $1 - o(1/\log n)$, among the first $O(\log \log n)$ column pairs along that row permutation some admits a legal toggling turn. We state these precisely, report the numerical evidence (up to $n = 150$, including, through a sampler of completions of a fixed rectangle whose connectivity and mixing are not proved, types of exponentially small probability), and locate the remaining difficulty: every exact identity we have found inside the conditioned space is fair, upper bounds on events at sparse sets of cells in a uniform completion of a fixed $2 \times n$ rectangle are now available, lower bounds at scale n on cycle lengths at the first square-selected column pair and the fixed-position form of the needed “different cycles” bound are proved, and what is missing is a lower bound for a constant-probability event at positions selected by the square — at the first candidate, a positive density of good pairs across the cut between the two marked rows. Along the way we observe that the published proof of the upper bound in Lemma 3.12 of Cavenagh–Greenhill–Wanless has a gap for every split other than $(2, 2)$; our results do not depend on that bound, and we re-derive from the surviving direction the tail estimate on the number of cycles that the reduction theorem needs.

1 Introduction

Let P_n be the law of the cycle type of $\sigma = \sigma_{1,2}(L)$ for a uniform Latin square L of order n , and Q_n the law of the cycle type of a uniform derangement of $[n]$. Cavenagh, Greenhill and Wanless [3, Conj. 6.1] conjectured:

Conjecture 1.1. $d_{\text{TV}}(P_n, Q_n) = o(1)$ as $n \rightarrow \infty$.

For $\lambda \vdash n$ with all parts ≥ 2 let $C_n(\lambda)$ be the number of completions of a fixed $2 \times n$ Latin rectangle of cycle type λ to a Latin square, divided by $(n-2)!$ (the count does not depend on the rectangle chosen, since rectangles of the same type are isotopic by a row-preserving isotopy), and let $\gamma(\lambda)$ be the number of permutations of type λ . Then $P_n(\lambda) \propto \gamma(\lambda)C_n(\lambda)$ and $Q_n(\lambda) \propto \gamma(\lambda)$, so the conjecture says that C_n is nearly constant on the types that carry most of the mass. CGW proved $C_n(\lambda)/C_n(\mu) \in [\frac{1}{2}, \frac{3}{2}]$ for adjacent types (one part of λ split into two parts of μ , both ≥ 2), and used it to show that the two laws agree within a factor $n^{3/2}$ on every type with at most $\frac{6}{5} \log n$ parts; the upper bound $\frac{3}{2}$, and with it the upper half of the latter statement, is affected by the gap described in Remark 3.5 and [6]. One marginal of the conjecture is known: by Kwan, Petrova and Sawhney’s law of large numbers for the number of odd rows [9, Thm. 1.3(1)] and the exchangeability of rows, $P_n(\sigma \text{ even}) = \frac{1}{2} + o(1)$, as for uniform derangements (§9).

Status of the arguments. The proofs below have been checked by automated referees (independent language-model agents with access to the sources and the data) and, where feasible, by exhaustive computation at $n \leq 7$ and by sampling (this includes §6(e), Remark 3.4, §7 and §8, added last; the latter two by two independent automated referees each, and §8 also by an independent cold read); no human referee has read them. The literature statements of §9 were checked against the cited papers by the same means.

Results. Throughout, a *mark* is a column pair $P = \{j, j'\}$ with $j' = \sigma^\alpha(j)$ on an m -cycle of σ , $m = \alpha + \beta$, $\alpha, \beta \geq 2$; $X = X(n, \lambda; \alpha, \beta)$ is the set of marked squares of type λ ; the *frame* of P is the permutation of rows $\pi(r) = r'$ with $L(r', j') = L(r, j)$; and P is a *B-pair* if rows 1, 2 lie in one cycle of π , an *A-pair* otherwise. Our results are:

- (1) *Marked-pair identity* (Theorem 2.4): $2C_n(\lambda)/C_n(\mu) - 1 = \mathbb{P}_X[B]/\mathbb{P}_X[A]$ exactly, for every $n, \lambda, \alpha, \beta$.
- (2) *Reduction* (Theorem 3.3): if $|\mathbb{P}_X[B]/\mathbb{P}_X[A] - 1| \leq \delta(n)$ for all adjacent pairs with $\alpha \neq \beta$ in which μ has at most $\lceil 16 \log n \rceil$ parts (equivalently $|\mathbb{P}_X[B] - \frac{1}{2}| \leq \frac{1}{4}\delta(1 + o(1))$), then $d_{\text{TV}}(P_n, Q_n) = O(\delta \log n + n^{-1+o(1)})$. The tail estimate it needs on the number of cycles of σ is proved from the trivial half of (1) alone (Proposition 3.2).
- (3) *Ladder identity* (Theorem 4.2): with the ladder pairs $(x_k, y_k) = (\pi^{-k}(1), \pi^{-k}(2))$, $1 \leq k < K_0$, and “flippable” meaning that j, j' lie in different cycles of the permutation ρ_{x_k, y_k} of columns induced by rows x_k, y_k ,

$$\mathbb{P}_X[B] - \mathbb{P}_X[A] = \mathbb{P}_X[B, \text{Flip} = \emptyset] - \mathbb{P}_X[A, \text{Flip} = \emptyset], \quad |\mathbb{P}_X[B] - \frac{1}{2}| \leq \frac{1}{2}\mathbb{P}_X[\text{Flip} = \emptyset].$$

So Conjecture 1.1 follows from (L): $\mathbb{P}_X[\text{Flip} = \emptyset] = o(1/\log n)$.

- (4) *Witness lemma* (Proposition 5.2): (L), hence the conjecture, follows from Hypothesis 5.1: with probability $1 - o(1/\log n)$, a long ladder contains a pair (x_k, y_k) whose ρ_{x_k, y_k} -cycle through j has length $\leq \log^5 n$ and avoids j' (such a pair is flippable outright), and, for the offset diagonals of short-arc instances, that a constant fraction of the expected number of diagonals carry one. This is the reference hypothesis of the paper.
- (5) *Orbit method* (Propositions 6.3 and 6.6): the flippabilities of a family of ladder pairs are conditionally independent given the orbit of a group of commuting column-cycle turns. With a fixed matching of columns as coordinates this gives $\mathbb{P}_X[\text{Flip}_I = \emptyset] \leq \mathbb{E}_X \exp(-\frac{1}{2} \sum_{k \in I} \bar{y}_k)$ with an explicit “orbit weight” $\bar{y}_k$; with the *adaptive* coordinates $(\rho_k^t(j), \rho_k^t(j'))$ — the t -th columns along the permutation $\rho_k = \rho_{x_k, y_k}$ of the k -th ladder pair from the marked columns — every coordinate is separated by the mark, so its turn toggles whenever it is legal and clean, and $\mathbb{P}_X[\text{Flip}_I = \emptyset] \leq \mathbb{E}_X[2^{-|U|}]$, where U is the set of ladder pairs having a short clean candidate cycle (subject to a collision condition). These give two alternative sufficient conditions (Hypotheses 6.4 and 6.7), for which we see no implication to or from Hypothesis 5.1: the first needs a genericity constant that the switchings available inside X cannot provide, the second needs $C \log \log n$ rare local events on a long ladder where the witness lemma needs one.
- (6) *The first ladder pair alone* (Propositions 6.10 and 6.11): at the first pair (x_1, y_1) the adaptive turn needs no side conditions, and combined with the first ladder trade it gives

$$\mathbb{P}_X[B] - \mathbb{P}_X[A] = \mathbb{E}_X[(\mathbf{1}_B - \mathbf{1}_A)(\mathbf{1}_{F^c} - \mathbf{1}_F); G^c], \quad |\mathbb{P}_X[B] - \tfrac{1}{2}| \leq \tfrac{1}{2} \mathbb{P}_X[G^c],$$

where F is the flippability of (x_1, y_1) and G the event that one of the $\nu - 1$ column pairs $\{\rho^t(j), \rho^t(j')\}$ along $\rho = \rho_{x_1, y_1}$ (ν = the length of the shorter arc or cycle of ρ at the mark) induces a permutation of rows that separates x_1 from y_1 with a legal cycle. So the conjecture also follows from Hypothesis 6.12: $\sup_{\lambda, \alpha, \beta} \mathbb{P}_X[G^c] = o(1/\log n)$. With $M = C \log \log n$, $\mathbb{P}_X[G^c] \leq \mathbb{P}_X[\nu \leq M] + \mathbb{P}_X[\nu > M, \text{ the first } M \text{ candidates bad}]$, and the second term is $o(1/\log n)$ for C large (depending on c) as soon as the conditional probability that candidate $t \leq M$ is bad, given $\nu > t$ and that its predecessors are bad, is at most $1 - c$; so the condition is an upper bound on a short arc at the frame-selected pair together with a constant conditional hazard over $O(\log \log n)$ events, in place of a long ladder of rare events, and it is the one hypothesis whose quantity the data measure directly ($n\mathbb{P}_X[G^c]$ is 5.5 for $30 \leq n \leq 150$, a first-moment fact that says nothing about the rate). Again no implication to or from Hypothesis 5.1 is known.

- (7) *The short-arc term is proved* (§7): the per-cell bound $\mathbb{P}[L(r, c) = s \mid L \supseteq P] \leq \Delta/n$ of Allsop and Morris [1], and with it the fixed-cell bound $\mathbb{P}[Q \subseteq L \mid L \supseteq P] \leq (\Delta/n)^{|Q|}$ for sparse Q , hold when P contains two complete rows (Theorem 7.1, Corollary 7.2; Δ absolute, ≤ 66), by running their switching proof with every column-cycle and cross-switch kept away from those rows. Hence $\mathbb{P}_X[\nu \leq M] \leq 8\Delta^{2M+2}/n$ for every type and mark (Corollary 7.3), and the conjecture follows from the *hazard hypothesis* (Hypothesis 7.4): with probability $1 - o(1/\log n)$, among the first $C \log \log n$ candidates some is good — which holds if each candidate is good with conditional probability $\geq c$ given that the earlier ones are bad (in the data a candidate is good with conditional probability ≈ 0.4).
- (8) *The hazard term: identity, lengths, fixed positions* (§8): the conditional hazard at each candidate equals a ratio of two local move counts, one in each of two conditioned ensembles (Proposition 8.1, exact); at the first candidate, conditionally on x_1, y_1 sharing a cycle

there, that cycle has length $\geq \delta n$ and its two arcs length $\geq \delta^3 n$ with probability $1 - O(\delta)$ (Proposition 8.3), by a switching of the Allsop–Morris type with the inverse pinned by the two marked rows; and for two *fixed* free rows and a *fixed* column pair, in the completions of any $2 \times n$ rectangle, the rows lie on different cycles with probability $\geq 1/30 - 3/n$ (Proposition 8.5), by row exchangeability on top of the same switching. What the first candidate still needs, and what none of these give, is a positive density of *good* pairs across the x_1 – y_1 cut without the exchangeability of y_1 ; we explain why the cross-switch counting of [1] does not supply it.

All these hypotheses are of one logical type: a lower bound on the probability of a configuration at rows or columns selected by the frame of the mark, inside the space conditioned on the type of rows 1, 2; by Remark 3.4 each is needed only for types with at most $\lceil 16 \log n \rceil$ parts. Their heuristic margins are large; none of the exact tools of this paper gives a lower bound of that kind. Upper bounds, for events at a sparse set of cells, are available after §7; the fixed-cell bounds in the literature for Latin squares and rectangles stop short of the conditioning on two full rows, and §7 extends the most recent of them, that of Allsop and Morris [1], to that case (§9 records the rest). The data (Jacobson–Matthews samples, and a sampler of completions of a fixed $2 \times n$ rectangle that reaches types of P_n -probability $e^{-\Theta(n)}$, Remark 6.14; its connectivity and mixing are not proved) give point estimates of $\mathbb{P}_X[B]$ within 0.003 of $\frac{1}{2}$ in every tested class at $n \leq 50$ and witness events that are independent to three digits along the ladder; they do not reach the regime of any of the hypotheses, and the empirical sizes quoted for the hypothesised quantities are first-moment information only.

A gap in CGW’s Lemma 3.12. Case 3 of the splitting procedure in [3] (cross-switch at $\{\omega j, \omega j'\}$, then backflip at $\{j, j'\}$) does not behave as claimed: the cross-switch reverses the ω -arc from ωj to $\omega j'$ (or the complementary arc, depending on the smallest corner symbol), which contains j' (resp. j), so $\{j, j'\}$ ends at distance 2. For $\min(\alpha, \beta) \geq 3$ the backflip then lands in the wrong type; for $\min(\alpha, \beta) = 2 < \max(\alpha, \beta)$ the type is right but the edge is not an edge of the joining multigraph, so the identity $G_S = G_J$ on which the double count rests fails (already at $n = 5$, where we computed both multigraphs exactly). We give the details in a separate note [6]; here we only record (Remark 3.5) which published statements depend on it. Nothing in the present paper does: the lower bound $\frac{1}{2}$ of CGW’s lemma is the trivial half of Theorem 2.4, and the upper bound $\frac{3}{2}$ is, by the same theorem, $\mathbb{P}_X[B] \leq \frac{2}{3}$ — a weak form of what we are after. By Theorem 4.2 it would follow from $\mathbb{P}_X[\text{Flip} = \emptyset] \leq \frac{1}{3}$.

Conventions. $L(r, c)$ is the entry in row r , column c . For rows x, y , $\rho_{x,y}$ is the permutation of columns with $L(y, \rho_{x,y}(q)) = L(x, q)$, so that $\sigma = \rho_{1,2}$ up to inversion; a *row cycle* of (x, y) is a cycle of $\rho_{x,y}$. For columns q, c , $\delta_{q,c}$ is the permutation of rows with $L(\delta(r), c) = L(r, q)$, and its cycles are the *column cycles* of (q, c) . A *trade* of rows x, y along a row cycle swaps their entries on the columns of that cycle; a *turn* of a column cycle swaps the entries of columns q, c in its rows. Both are Latin trades (the result is again a Latin square). For $\lambda \vdash n$, $S(\lambda)$ is the set of Latin squares of order n with first row the identity whose rows 1, 2 have cycle type λ , so $|S(\lambda)| = \gamma(\lambda)(n-2)!C_n(\lambda)$; we also write $D_\lambda = \gamma(\lambda)$ for the number of derangements of type λ . A column pair $\{j, j'\}$ is *admissible* if $j' \notin \{j, \sigma(j), \sigma^{-1}(j)\}$ (CGW’s condition for the pair operations; every marked pair is admissible since $\alpha, \beta \geq 2$). Following CGW, the *flip* at an A -pair whose columns lie in different σ -cycles, and the *backflip* (or *unflip*) at an A -pair whose columns lie in one σ -cycle, is the turn of the column cycle of the pair through row 2 (it joins,

resp. splits, the σ -cycles involved); the *switch* at a pair whose columns lie in different σ -cycles is the trade of rows 1, 2 along the row cycle through one of the two columns, selected by CGW's rule (the cycle whose columns carry the smaller minimum symbol). We call Conjecture 1.1, which asks only for convergence in total variation, the *weak* conjecture; the stronger statements suggested by Cameron ("tends very rapidly") and by the data (pointwise ratios $1 + o(1)$) are not addressed here. Jacobson–Matthews sampling [8] is used for all numerical data; the scripts and logs are in the project repository.

2 The marked-pair identity

Throughout, squares are normalised to have first row id; CGW's column map ω , defined by $L(1, \omega(j)) = L(2, j)$, is then $\omega(j) = L(2, j)$, and we take $\sigma = \omega$ (only the cycle type matters, and it is the same for ω and ω^{-1}). For the proofs it is convenient to work with *unnormalised* squares (equivalently, symbol-relabelling classes): every instance count below then acquires a common factor $n!$, which cancels from every ratio, and the maps we use (flip, switch) act on unnormalised squares without any renormalisation step. Statuses, cycle types, and column positions are relabelling-invariant, so we freely state results in the normalised picture.

Definition 2.1 (marked instances). Let $\lambda \vdash n$ contain a part $m = \alpha + \beta$ ($\alpha, \beta \geq 2$), and let μ be λ with one part m split into α, β . The *splitting instance space* is

$$X = X(n, \lambda; \alpha, \beta) = \{(L', P) : L' \in S(\lambda), P = \{j, \sigma^\alpha(j)\} \text{ with } j \text{ in an } m\text{-cycle of } \sigma(L')\},$$

unordered pairs counted once. Thus every L' carries exactly $\lambda_m m$ marked pairs when $\alpha \neq \beta$ and $\lambda_m \alpha$ when $\alpha = \beta$. Since $\alpha, \beta \geq 2$ every marked pair is admissible, and $X = X_A \sqcup X_B$ according to the A/B status of P in L' . The *joining instance space* is

$$J = J(n, \mu; \alpha, \beta) = \{(L, P) : L \in S(\mu), P = \{j, j'\}, j \text{ in an } \alpha\text{-cycle, } j' \text{ in a different } \beta\text{-cycle}\},$$

again with $J = J_A \sqcup J_B$ by the status of P in L .

The counting arithmetic is immediate: with D_ν the number of derangements of cycle type ν and $N(\nu) = C_n(\nu)(n-2)!$,

$$\frac{2|X|}{|J|} = \frac{2 D_\lambda N(\lambda) \lambda_m m}{D_\mu N(\mu) \mu_\alpha \mu_\beta \alpha \beta} = \frac{2 C_n(\lambda)}{C_n(\mu)} \quad (\alpha \neq \beta), \quad (1)$$

using $D_\lambda/D_\mu = \alpha \mu_\alpha \beta \mu_\beta / (m \lambda_m)$; for $\alpha = \beta$ one has $D_\lambda/D_\mu = \alpha^2 \mu_\alpha (\mu_\alpha - 1) / (m \lambda_m)$, $|J| = |S(\mu)| \binom{\mu_\alpha}{2} \alpha^2$ and $|X| = |S(\lambda)| \lambda_m \alpha$, and (1) holds verbatim. The content of the identity is therefore the following count.

Lemma 2.2 (cycle trades preserve pair status). *Let $P = \{j, j'\}$ be an admissible column pair of a Latin square L and let C be one column cycle of P . Trading C (swapping the entries of columns j, j' in the rows of C) inverts the pair permutation π_P on C and fixes it elsewhere; the partition of rows into column cycles of P is unchanged. In particular the A/B status of P itself is invariant under the CGW flip at P , which trades the full column cycle through row 2.*

Proof. Write $\pi = \pi_P$, so $\pi(r)$ is the row of column j' holding the symbol $L(r, j)$; C is π -closed and the trade swaps $L(r, j) \leftrightarrow L(r, j')$ for $r \in C$. Fix $r \in C$ and let π' be the new pair permutation.

The new $L(r, j)$ is the old $L(r, j')$; the row of the new column j' holding that symbol is found by solving (new $L(s, j')$) = (old $L(r, j')$): for $s \in C$ the new $L(s, j')$ is the old $L(s, j)$, and old $L(s, j) = \text{old } L(r, j')$ exactly when $\pi(s) = r$. Hence $\pi' = \pi^{-1}$ on C ; off C nothing changes. A cycle inverted in place has the same support, so the partition into cycles is unchanged. $\square$

Lemma 2.3 (switch toggle). *Let $(L, P) \in J$ and let Q be the row cycle through one column of P (the two row cycles are distinct), selected by CGW's min-symbol rule. Trading Q (the CGW switch) changes π_P by precomposition with the transposition of rows 1, 2: the cycles of π_P through rows 1, 2 are split if rows 1, 2 are in one cycle, and joined if in two. Hence the switch toggles the status of P , preserves $S(\mu)$ and the instance data (it inverts one σ -cycle in place), and is an involution: the trade preserves the traded cycle's symbol set (the two rows carry the same symbols on its columns), so the min-symbol rule re-selects the same cycle. Consequently $|J_A| = |J_B|$, exactly.*

Proof. Say Q is the row cycle through j ; its columns include j but not j' . The trade swaps the row-1 and row-2 entries in the columns of Q ; for π_P only column j matters, where the entries of rows 1, 2 are exchanged. Thus $\pi'_P = \pi_P \circ \tau_{12}$, and multiplying by a transposition splits or joins the cycles through its two points. Rows 1, 2 lie in one cycle of π_P exactly when P is a B -pair. The trade swaps the two rows' entries along the columns of Q , which inverts the corresponding σ -cycle in place: the cycle type and the column sets of all row cycles are preserved, so $(L, P) \mapsto (\text{sw}(L), P)$ maps J to J . $\square$

Theorem 2.4 (marked-pair identity). *For every n , λ , and split $m = \alpha + \beta$ with $\alpha, \beta \geq 2$ (including $\alpha = \beta$),*

$$|J_A| = |J_B| = |X_A|, \quad \text{hence} \quad \boxed{\frac{C_n(\lambda)}{C_n(\mu)} = \frac{|X|}{2|X_A|}, \quad 2 \frac{C_n(\lambda)}{C_n(\mu)} - 1 = \frac{|X_B|}{|X_A|}.}$$

In particular $C_n(\lambda) > \frac{1}{2} C_n(\mu)$ unconditionally, and Hypothesis 3.1 (EQ(δ)) is equivalent to

$$\left| \mathbb{P}(P \text{ is a } B\text{-pair}) - \frac{1}{2} \right| \leq \frac{\delta}{4}(1 + o(1)),$$

for (L', P) uniform on X : the law of a single bit — whether rows 1, 2 lie in one column cycle of the marked pair $\{j, \sigma^\alpha(j)\}$ — of a single uniform square of $S(\lambda)$.

Proof. By (1) it suffices to prove $|J| = 2|X_A|$. The flip at an A -pair P whose columns lie in one row cycle at σ -distances (α, β) splits that cycle into an α - and a β -cycle and maps $S(\lambda) \rightarrow S(\mu)$; by Lemma 2.2 the pair P is still an A -pair afterwards, so $x \mapsto \text{unflip}(x)$ maps X_A into J_A . The flip at a fixed pair is an involution (the traded cycle again contains row 2, with the same row support), so $\text{join} = \text{flip}$ inverts it: $X_A \leftrightarrow J_A$ is a bijection. Composing with the switch of Lemma 2.3 gives a bijection $X_A \leftrightarrow J_B$, namely $x \mapsto \text{sw}(\text{unflip}(x))$ with inverse $(L, P) \mapsto \text{flip}(\text{sw}(L))$ — the map “switch-then-flip”, which is exactly the CGW joining step at a B -pair. Hence $|J| = |J_A| + |J_B| = 2|X_A|$. $\square$

3 The reduction theorem

Hypothesis 3.1 (B-pair equidistribution, EQ(δ)). *For every $\lambda \vdash n$ with a part $m = \alpha + \beta$, $\alpha \neq \beta$, $\alpha, \beta \geq 2$, and μ the split type, $|\bar{q}(\mu \rightarrow \lambda) - 1| \leq \delta(n)$, where $\bar{q}(\mu \rightarrow \lambda) := 2C_n(\lambda)/C_n(\mu) - 1 = \mathbb{P}_X[B]/\mathbb{P}_X[A]$ by Theorem 2.4.*

By Theorem 2.4, $\text{EQ}(\delta)$ is the statement $|\mathbb{P}_X[B] - \frac{1}{2}| \leq \frac{\delta}{4}(1 + o(1))$ for all adjacent pairs with $\alpha \neq \beta$. The proof below needs a tail bound on the number κ of cycles of σ ; we first prove it from the trivial half of Theorem 2.4 alone, so that nothing depends on the upper bound of [3, Lemma 3.12] (Remark 3.5).

Proposition 3.2 (the number of cycles, from the surviving direction only). *Let L be a uniform Latin square of order n , σ the column permutation of rows 1, 2, N_ℓ its number of ℓ -cycles and $\kappa = \sum_\ell N_\ell$. Then, using only the identity $|J| = 2|X_A|$ of Theorem 2.4 and $|X_A| \leq |X|$:*

(i) $\mathbb{E}[N_\ell] \leq 2n/(n - \ell)$ for every $\ell < n$; in particular $\mathbb{E}[N_\ell] \leq 4/\ell$ for $\ell \leq n/2$.

(ii) For every $k \geq 1$, with $(\kappa)_k = \kappa(\kappa - 1) \cdots (\kappa - k + 1)$,

$$\mathbb{E}[(\kappa)_k] \leq 2^{k-1} \sum_{\ell_1, \dots, \ell_k \geq 2} \frac{\mathbb{E}[mN_m]}{\ell_1 \cdots \ell_k} \Big|_{m=\ell_1+\dots+\ell_k} \leq 2^{k-1} \left(4 + \frac{k^2}{\log n}\right) (\log n)^k.$$

(iii) For $A \geq 3$, $\mathbb{P}[\kappa \geq A \log n] \leq (7 + \log n) n^{-\log((A-1)/2)}$, which is $O(n^{-2} \log n)$ for $A \geq 16$; and $\mathbb{P}[\text{all cycles of } \sigma \text{ have the same length } \ell < n] \leq 4d(n)/n$, $d(n)$ the number of divisors.

Proof. (i) For $\beta \neq \ell$ the joining instances $(L, \{j, j'\})$ with j in an ℓ -cycle and j' in a β -cycle number $\ell N_\ell \beta N_\beta$ per square, the splitting instances (pairs at distance ℓ on an $(\ell + \beta)$ -cycle) number $(\ell + \beta)N_{\ell+\beta}$, and $|J| = 2|X_A| \leq 2|X|$ gives $\mathbb{E}[\ell N_\ell \beta N_\beta] \leq 2\mathbb{E}[(\ell + \beta)N_{\ell+\beta}]$; for $\beta = \ell$ the same count with $\ell^2 \binom{N_\ell}{2}$ and $\ell N_{2\ell}$ gives $\mathbb{E}[\ell^2 N_\ell (N_\ell - 1)] \leq 4\ell \mathbb{E}[N_{2\ell}]$. Summing the first over $\beta \neq \ell$ ($\sum_\beta \beta N_\beta = n$) and adding the second (its right side is the missing summand $m = 2\ell$): $\mathbb{E}[\ell N_\ell (n - \ell N_\ell)] + \mathbb{E}[\ell^2 N_\ell (N_\ell - 1)] \leq 2\mathbb{E}[\sum_{m \geq \ell+2} m N_m] \leq 2n$. Hence $n\mathbb{E}[N_\ell] = \mathbb{E}[\ell N_\ell (n - \ell N_\ell)] + \mathbb{E}[\ell^2 N_\ell (N_\ell - 1)] + \ell \mathbb{E}[N_\ell] \leq 2n + \ell \mathbb{E}[N_\ell]$.

(ii) Count the instances $(L, (C_1, j_1), \dots, (C_k, j_k))$: C_i pairwise distinct σ -cycles of L of lengths ℓ_i and $j_i \in C_i$; there are $\ell_1 \cdots \ell_k$ of them per ordered k -tuple of distinct cycles with these lengths. Join C_1 and C_2 at the pair $\{j_1, j_2\}$ by a joining step: flip if the pair is an A -pair; if it is a B -pair, first trade the row cycle C_2 of j_2 (a switch in the sense of Lemma 2.3, whose proof applies to the row cycle through either column of the pair; it inverts C_2 in place and does not touch the cycle of j_1), then flip. The flip replaces σ by $\sigma \circ (j_1 j_2)$ on rows 1, 2, so the result has j_1, j_2 on one cycle of length $\ell_1 + \ell_2$ with $j_2 = \sigma^{\ell_2}(j_1)$ whether or not the trade was applied (CGW Lemma 3.3), every other cycle of σ (in particular $C_3, \dots, C_k$) untouched, and the map is at most two-to-one (the backflip at $\{j_1, j_2\}$ returns L or the traded square). Repeat with $\{j_1, j_3\}, \dots, \{j_1, j_k\}$, always trading the cycle of the new column when needed, so that the cycle containing j_1 is never inverted: after step i every earlier j_r is shifted by exactly ℓ_{i+1} . After $k - 1$ steps all j_i lie on one cycle of length $m = \sum \ell_i$ at the positions $j_i = \sigma^{\ell_i + \dots + \ell_k}(j_1)$, determined by j_1 and $(\ell_1, \dots, \ell_k)$, and the composite is at most 2^{k-1} -to-one. (With CGW's min-symbol rule the trade may invert the merged cycle and the positions are not determined; the count would then be 2^{2k-3} -to-one.) Hence the number of instances with lengths (ℓ_i) is at most $2^{k-1} \sum_{L'} m N_m(L')$, i.e. $\mathbb{E}[\ell_1 \cdots \ell_k T(\ell_1, \dots, \ell_k)] \leq 2^{k-1} \mathbb{E}[m N_m]$ with T the number of ordered k -tuples of distinct cycles with these lengths; summing over (ℓ_i) gives the first bound, since $\sum_{(\ell_i)} T = (\kappa)_k$. For the second, split at $m \leq n/2$, where $\mathbb{E}[m N_m] \leq 4$ by (i) and $\sum_{m \leq n/2} \sum_{\sum \ell_i = m} \prod \ell_i^{-1} \leq (\sum_{\ell \leq n/2} \ell^{-1})^k \leq (\log n)^k$, and $m > n/2$, where some $\ell_i \geq m/k$, so $\sum_{\sum \ell_i = m} \prod \ell_i^{-1} \leq k \cdot (k/m)(\log n)^{k-1}$, and $\sum_{m > n/2} m N_m \cdot k^2 (\log n)^{k-1} / m \leq k^2 (\log n)^{k-1}$ because at most one cycle is longer than $n/2$.

(iii) Markov for $(\kappa)_k$ with $k = \lceil \log n \rceil \leq \log n + 1$: on $\{\kappa \geq K\}$, $(\kappa)_k \geq (K - k + 1)^k \geq ((A - 1) \log n)^k$ for $K = A \log n$, and $k^2 / \log n \leq \log n + 3$, so $\mathbb{P}[\kappa \geq K] \leq 2^{k-1} (4 + k^2 / \log n) (\log n)^k / ((A - 1) \log n)^k \leq (7 + \log n) (2 / (A - 1))^k \leq (7 + \log n) (2 / (A - 1))^{\log n}$ for $A \geq 3$, and $(A - 1) / 2 \geq e^2$ for $A \geq 16$. For the equal-length types, $\ell \mid n$ and $\ell < n$ force $\ell \leq n/2$, so $\mathbb{P}[N_\ell = n/\ell] \leq \mathbb{E}[N_\ell] \ell / n \leq 4/n$; $\ell = n$ is the n -cycle. $\square$

Theorem 3.3. Assume $\text{EQ}(\delta)$ with $\delta = \delta(n) \rightarrow 0$, or only its instances in which μ has at most $\lceil 16 \log n \rceil$ parts. Then

$$d_{\text{TV}}(P_n, Q_n) = O(\delta \log n + n^{-1+o(1)}).$$

In particular $\text{EQ}(o(1/\log n))$ implies Conjecture 1.1.

Proof. Call λ plain if not all its parts are equal (so the only non-plain types are (ℓ^κ) , $\kappa \geq 2$).

Routing. Every plain λ with $\kappa = \kappa(\lambda)$ parts is reachable from (n) by a chain of $\kappa - 1$ splits, each with $\alpha \neq \beta$. Build the chain in reverse, splitting parts of λ off a shrinking residual, in a suitable order; a collision means splitting a part p off a residual R with $p = R - p$. If the two largest parts of λ are unequal, split smallest-first and let them be the final pair: while at least three parts remain, the split-off minimum p satisfies $R - p \geq 2p > p$ (all parts are ≥ 2 and at least two other parts remain), so no collision occurs, and the final split is unequal by hypothesis. If instead the two largest parts are equal, say to $\ell_{\max}$, then $\ell_{\max} \neq n/2$ (else $\lambda = (n/2, n/2)$ is not plain), so we may split one copy of $\ell_{\max}$ off first without collision, then proceed smallest-first among the remaining parts, reserving for the final split the pair $(b, \ell_{\max})$ where $b < \ell_{\max}$ is the smallest part: at every intermediate step the residual minus the split-off minimum still contains the reserved $\ell_{\max}$ and the reserved b , hence exceeds the minimum (also when a further copy of $\ell_{\max}$ is the split-off minimum: $R - p \geq \ell_{\max} + b > p$), and the final split is unequal.

For each step of such a chain, Theorem 2.4 and $\text{EQ}(\delta)$ give $|S(p_{i+1})|/|S(p_i)| = (\gamma(p_{i+1})/\gamma(p_i))(1 + \varepsilon_i)$ with $|\varepsilon_i| \leq \delta$. Multiplying along the chain,

$$\frac{P_n(\lambda)}{P_n((n))} = \frac{\gamma(\lambda)}{\gamma((n))} \prod_{i=1}^{\kappa-1} (1 + \varepsilon_i)^{\pm 1}, \quad |\varepsilon_i| \leq \delta,$$

so for all plain λ with $\kappa(\lambda) \leq K$, $P_n(\lambda) = Q_n(\lambda) Z^{-1} (1 + O(K\delta))$ (the bound is vacuous unless $K\delta = O(1)$, which is the case of interest) where $Z > 0$ is a normalising constant independent of λ (namely $Z = \sum_{\nu} \gamma(\nu) C_n(\nu) / (|\mathcal{D}_n| C_n((n)))$).

Let $K = \lceil A \log n \rceil$ with $A = 16$, and let $\mathcal{B}$ be the set of types that are either non-plain or have $\kappa > K$. Under Q_n , Markov's inequality with the generating function $\sum_{\pi \in \mathcal{D}_n} 2^{\kappa(\pi)} = n! [x^n] e^{-2x} (1-x)^{-2} \leq n \cdot n!$ (cf. [3, proof of Lemma 4.4]) gives $Q_n(\kappa > K) \leq e n 2^{-K} = O(n^{-2})$ for $A \log 2 \geq 3$; also $Q_n(\text{non-plain}) \leq \sum_{\ell \mid n, \ell < n} \gamma((\ell^{n/\ell})) / D_n = O(n^{-1})$. Under P_n , Proposition 3.2 gives $P_n(\kappa > K) = O(n^{-2} \log n)$ for $A \geq 16$ and $P_n(\text{non-plain}) = O(n^{-1+o(1)})$ without any input from [3] beyond the identity of Theorem 2.4.

On the complement of $\mathcal{B}$, summing $P_n = Q_n Z^{-1} (1 + O(K\delta))$ over $\lambda \notin \mathcal{B}$ and using the tail bounds shows $Z^{-1} = 1 + O(K\delta) + O(n^{-1+o(1)})$, whence $\sum_{\lambda \notin \mathcal{B}} |P_n - Q_n| = O(K\delta) + O(n^{-1+o(1)})$. Combining with the tails, $d_{\text{TV}}(P_n, Q_n) = O(\delta \log n + n^{-1+o(1)})$. $\square$

Remark 3.4 (which types are needed). The proof uses $\text{EQ}(\delta)$ only along the chains from (n) to the plain types with at most $K = \lceil 16 \log n \rceil$ parts, i.e. for adjacent pairs $\mu \rightarrow \lambda$ with $\kappa(\mu) \leq K$ (and $\alpha \neq \beta$; every intermediate type of a chain is plain, since the residual part always exceeds every part already split off, so λ and μ are both plain). Consequently each of the sufficient

conditions below (Hypotheses 5.1, 6.4, 6.7, 6.12), which are stated as suprema over all λ, α, β , is needed only for plain types λ with fewer than K parts and plain splits μ . Asymptotically this excludes the types with $\Theta(n)$ parts, such as the families $(4, 2^{(n-4)/2})$ of Remark 6.14 (at the sizes we sample, $n \leq 150$, every tested type still has fewer than $\lceil 16 \log n \rceil$ parts; the pairs $(4, 2^{13}) \rightarrow (2^{15})$ and $(6, 3^8) \rightarrow (3^{10})$ there are not needed because $\alpha = \beta$ and μ is not plain); it does not remove the conditioning itself, which is on two full rows whatever the type.

Remark 3.5 (the gap in CGW’s Lemma 3.12). In [3], the direction $|S(\lambda, F)|N_\lambda \geq \frac{1}{2}|S(\mu, F)|N_\mu$ of Lemma 3.12 uses only that each splitting pair carries at most two edges and each joining pair at least one; by Theorem 2.4 it is the trivial inequality $|X| \geq |X_A|$. The direction $\leq \frac{3}{2}$ requires every splitting pair of every λ -square to carry exactly two edges that are also edges of the joining multigraph, and case 3 of the splitting procedure fails to supply them for every split other than $(2, 2)$ [6]. Consequently the upper bound of CGW’s Theorem 3.13, the lower bound of their Lemma 4.1, the parity bounds of Theorem 4.3, the bound $\mathbb{P}[\sigma \text{ is an } n\text{-cycle}] \leq 2n^{-2/3}$ of Lemma 4.4, both halves of Corollary 4.5, and the results of their Section 4.2 that rest on these (Corollary 4.6, Lemma 4.7, Theorems 4.9 and 4.11, Corollary 4.10 — the last superseded by [10]) are not established by the published argument. The *statements* of Theorem 4.9 (at most $9\sqrt{n}$ cycles with probability $1 - o(e^{-\sqrt{n}/4})$) and of the upper half of Corollary 4.5 (with $n/2$ in place of $n^{1/3}$) do follow from the surviving direction: Proposition 3.2(ii) with $k = \sqrt{n}$ gives $\mathbb{P}[\kappa \geq 9\sqrt{n}] \leq \exp(-(\frac{1}{2} - o(1))\sqrt{n} \log n)$, and $C_n(\lambda) \leq 2^{\kappa-1}C_n((n))$ gives $P_n(\lambda) \leq \frac{n}{2} \prod_i (2/c_i)$ for a type with parts c_i ; this repairs the use of both in [7, Thm. 9]. The surviving direction gives no upper bound on coarse types: in particular we know of no proof that two rows of a random Latin square form a single n -cycle with probability $o(1)$ (Lemma 4.4’s $2n^{-2/3}$; the second-moment method with [10] and Proposition 3.2(i) gives only $\leq \frac{5}{6} + o(1)$), and the parity theorem of [4] uses the unproved direction at odd splits (details in [6]); that theorem is expected to be true independently of the gap, being a special case of [9, Thm. 1.3(4)], whose written proof uses [4] as a black box but whose Remark 6.6 sketches an independent re-proof, which we have not checked. For $\min(\alpha, \beta) = 2$ the count can be corrected with the constant 2 in place of $\frac{3}{2}$ [6, Prop. 3]: for marks at distance 2 or $m - 2$ on an m -cycle, $C_n(\lambda)/C_n(\mu) \leq 2$, i.e. $\mathbb{P}_X[A] \geq \frac{1}{4}$ unconditionally. The present paper uses none of the affected statements of [3]: Proposition 3.2 replaces Corollary 4.5 in the proof of Theorem 3.3. The exact completion counts tabulated in [3] give $C_n(\lambda)/C_n(\mu) \in [0.985, 1.0035]$ for every adjacent pair at $7 \leq n \leq 11$ (and $= \frac{3}{2}$ at $n = 5$, $\frac{7}{8}, \frac{11}{14}, \frac{21}{22}$ at $n = 6$), so the statement itself is not in doubt.

4 The ladder identity

Notation: $(L, P) \in X = X(n, \lambda; \alpha, \beta)$, $P = \{j, j'\}$, frame $\pi = \pi_P$ ($L(\pi(r), j') = L(r, j)$), C_r the π -cycle of row r , $\rho_{x,y}$ the column permutation of rows x, y ($L(y, \rho_{x,y}(q)) = L(x, q)$), (x, y) *flippable* iff j, j' lie in different cycles of $\rho_{x,y}$, and the j' -trade at (x, y) exchanges rows x, y on the $\rho_{x,y}$ -cycle through j' , which replaces π by $(xy) \circ \pi$ if (x, y) is flippable and by $(xy) \circ \pi \circ (xy)$ otherwise. $B = \{1 \sim 2\}$ (rows 1, 2 in a common frame cycle), A its complement, $\mathbf{1}_B$ the bit; for B -instances d_{12} and d_{21} are the numbers of π -steps from 1 to 2 and from 2 to 1.

Definition 4.1 (ladder). Put $K_0 = \min(|C_1|, |C_2|)$ on A and $K_0 = \min(d_{12}, d_{21})$ on B , and for $1 \leq k < K_0$ let

$$x_k = \pi^{-k}(1), \quad y_k = \pi^{-k}(2), \quad T_k = \text{the } j'\text{-trade at } (x_k, y_k).$$

The pairs $\{x_k, y_k\}$ are the *ladder* of (L, P) , and $\text{Flip}(L, P) = \{k < K_0 : (x_k, y_k) \text{ flippable}\}$.

The rows $x_1, \dots, x_{K_0-1}$ are distinct rows of $C_1 \setminus \{1\}$ (on A) or of the arc from 2 to 1 (on B), the y_k likewise in $C_2 \setminus \{2\}$ or the arc from 1 to 2; so the $2(K_0 - 1)$ ladder rows are distinct and avoid rows 1, 2, and $K_0 \geq 2$: on A because π is a derangement, on B by the admissibility of the mark ($\pi(1) \neq 2, \pi(2) \neq 1$).

Theorem 4.2 (ladder identity). *For every n , λ and split $\alpha + \beta = m$ with $\alpha, \beta \geq 2$,*

$$\mathbb{P}_X[B] - \mathbb{P}_X[A] = \mathbb{P}_X[B, \text{Flip} = \emptyset] - \mathbb{P}_X[A, \text{Flip} = \emptyset], \quad \text{hence} \quad |\mathbb{P}_X[B] - \frac{1}{2}| \leq \frac{1}{2} \mathbb{P}_X[\text{Flip} = \emptyset].$$

The same identity holds with X replaced by the set of squares with a fixed pair of rows 1, 2 (of type λ) and a fixed mark P .

Proof. Fix $(L, P) \in X$ and write $K = \{1, \dots, K_0 - 1\}$. The ladder depends on (L, P) , so the T_k are not maps of X ; the statement to prove is that $(L, P) \sim T_S(L, P)$, $S \subseteq K$ (the trades at the unordered pairs of S , in any order), is an equivalence relation on X whose classes have exactly $2^{|K|}$ elements and on which the bit is $\mathbf{1}_B(T_S(L, P)) = \mathbf{1}_B(L, P) + |S \cap \text{Flip}| \bmod 2$. This follows once we know that each T_k is an involution of X which preserves K_0 , the set of unordered ladder pairs and the flippability of each of them, that T_k toggles the bit exactly when $k \in \text{Flip}$, and that trades at disjoint row pairs commute: then the same trades, at the same unordered pairs with the same flippabilities, are available from every point of a class (symmetry and transitivity), the $2^{|K|}$ points $T_S(L, P)$ are distinct (rows x_k, y_k are changed by T_k and by no other T_i), and the bit formula holds. If $\text{Flip} \neq \emptyset$ exactly half of the subsets S have $|S \cap \text{Flip}|$ odd, so B and A each contain half of the class; summing over classes gives the identity, and the bound follows from $\mathbb{P}[B] + \mathbb{P}[A] = 1$.

T_k is a well-defined involution of X . The j' -trade at (x_k, y_k) changes only rows $x_k, y_k \notin \{1, 2\}$, so σ and the mark are untouched and (L, P) stays in X ; it is an involution because the traded columns are again the ρ'_{x_k, y_k} -cycle through j' .

Effect on the frame. If $k \in \text{Flip}$ the new frame is $\pi' = (x_k y_k) \circ \pi$. On A this merges C_1 and C_2 : the new cycle runs $1 \rightarrow \pi(1) \rightarrow \dots \rightarrow x_{k+1} \rightarrow y_k \rightarrow \pi(y_k) \rightarrow \dots \rightarrow 2 \rightarrow \dots \rightarrow y_{k+1} \rightarrow x_k \rightarrow \dots \rightarrow 1$, so $d'_{12} = |C_1|$ and $d'_{21} = |C_2|$ (reading $x_{|C_1|} = 1, y_{|C_2|} = 2$). On B , x_k lies on the arc from 2 to 1 and y_k on the arc from 1 to 2 (as $k < \min(d_{12}, d_{21})$), and the split into $(x_k, \pi(x_k), \dots, \pi^{-1}(y_k))$ and $(y_k, \dots, \pi^{-1}(x_k))$ puts 1 (k steps after x_k , fewer than the d_{12} steps from x_k to y_k) in the first cycle and 2 in the second: the new state is in A with $|C'_1| = d_{12}$, $|C'_2| = d_{21}$. In both cases K_0 is preserved and the bit toggles. If $k \notin \text{Flip}$ the new frame is $(x_k y_k) \circ \pi \circ (x_k y_k)$: the cycle partition is unchanged up to exchanging the labels x_k, y_k , so K_0 and the bit are preserved.

The ladder is preserved. Let $\pi' = (x_k y_k) \circ \pi$ (the conjugation case is similar and easier). Since $\pi'^{-1}(r) = \pi^{-1}((x_k y_k)(r))$, the preimage chains of 1 and 2 under π' agree with those under π up to depth k and are exchanged beyond it: $\pi'^{-i}(1) = x_i$ and $\pi'^{-i}(2) = y_i$ for $i \leq k$, while $\pi'^{-(k+1)}(1) = \pi^{-1}(y_k) = y_{k+1}$ and then $\pi'^{-i}(1) = y_i, \pi'^{-i}(2) = x_i$ for $k < i < K_0$. So the unordered pairs $\{x_i, y_i\}$, $i \in K$, are the ladder of $T_k(L, P)$. In the conjugation case the same computation with $(x_k y_k) \circ \pi \circ (x_k y_k)$ gives $\pi'^{-i}(1) = (x_k y_k)(x_i)$, which again permutes the pairs $\{x_i, y_i\}$ among themselves (only $i = k$ is affected, by exchanging x_k, y_k).

Flippability is preserved and the T_k commute. T_k changes rows x_k, y_k only, and the flippability of (x_i, y_i) , $i \neq k$, depends only on rows x_i, y_i , which are different rows. Likewise T_i and T_k act on disjoint pairs of rows, each along a column set determined by its own two rows, so $T_i T_k = T_k T_i$. (After T_k the pair $\{x_i, y_i\}$ may have changed its index, but the trade at an unordered pair does not depend on the index or on the order of the two rows.)

The statement for fixed R holds because every T_k fixes rows 1 and 2. $\square$

Corollary 4.3 (the conjecture reduces to the ladder). *Write $f(n, \lambda; \alpha, \beta) = \mathbb{P}_X[\text{Flip} = \emptyset]$. By Theorem 2.4, $\bar{q}(\mu \rightarrow \lambda) - 1 = (\mathbb{P}_X[B] - \mathbb{P}_X[A])/\mathbb{P}_X[A]$ and $\mathbb{P}_X[A] \geq (1 - f)/2$, so $|\bar{q} - 1| \leq 2f/(1 - f)$: EQ(δ) holds with $\delta(n) = \max_{\lambda, \alpha \neq \beta} 2f/(1 - f)$. If this $\delta(n) \rightarrow 0$, Theorem 3.3 gives*

$$d_{\text{TV}}(P_n, Q_n) = O\left(\log n \cdot \max_{\lambda, \alpha \neq \beta} f(n, \lambda; \alpha, \beta) + n^{-1+o(1)}\right).$$

In particular the weak CGW conjecture follows from

$$(L) \quad \mathbb{P}_X[\text{no ladder pair is flippable}] = o(1/\log n) \quad \text{uniformly in } \lambda, \alpha, \beta. \quad (2)$$

Remark 4.4 (data). On Jacobson–Matthews samples with a uniformly random mark (9000, 7200, 3000 ladder-pair instances at $n = 30, 50, 100$; runs/s13/f1/): $\mathbb{P}[\text{Flip} = \emptyset \mid K_0] = 2^{-(K_0-1)}$ within statistical error for $K_0 \leq 6$ and consistent with it beyond; $\mathbb{P}[\text{Flip} = \emptyset] = 0.125, 0.079, 0.041 \approx 4/n$; the ladder pairs are flippable with frequency 0.5002 (62082 pairs at $n = 50$); and $\mathbb{P}[B] - \mathbb{P}[A] = -0.022 \pm 0.011, +0.008 \pm 0.012, -0.030 \pm 0.018$. On completions of a fixed rectangle (Remark 6.14; $7 \cdot 10^5$ marked instances at $n = 30$) $\mathbb{P}[\text{Flip} = \emptyset \mid K_0 - 1 \geq K] = 2^{-K}$ within noise to $K = 12$ –15, in every tested type.

4.1 Offset ladders and short arcs

The ladder is one of a family. For an integer c (an *offset*) put

$$K_c = \begin{cases} \min(|C_1| - c, |C_2|) & \text{on } A, \\ \min(d_{21} - c, d_{12}) & \text{on } B, \end{cases} \quad D_c = \{(x_{c+i}, y_i) = (\pi^{-(c+i)}(1), \pi^{-i}(2)) : 1 \leq i < K_c\},$$

so D_0 is the ladder and for $c > 0$ the x -row is c steps deeper than the y -row (x -depths $c + 1, \dots, c + K_c - 1$ stay inside $C_1 \setminus \{1\}$, resp. the arc from 2 to 1; y -depths $1, \dots, K_c - 1$ inside $C_2 \setminus \{2\}$, resp. the arc from 1 to 2). We take $c \geq 0$; the families with the y -row deeper are obtained by exchanging the roles of rows 1 and 2 (and of C_1, C_2, d_{12}, d_{21}) throughout.

Theorem 4.5 (offset ladders). *For every offset $c \geq 0$, on the domain $\mathcal{D}_c = \{K_c \geq 2\}$ (where $\mathcal{D}_c \neq \emptyset$) the j' -trades at the pairs of D_c are commuting involutions that preserve K_c , the unordered pairs of D_c and their flippabilities, and toggle the bit iff the traded pair is flippable. Hence, with Flip_c the set of flippable pairs of D_c ,*

$$\mathbb{P}[B, \mathcal{D}_c] - \mathbb{P}[A, \mathcal{D}_c] = \mathbb{P}[B, \mathcal{D}_c, \text{Flip}_c = \emptyset] - \mathbb{P}[A, \mathcal{D}_c, \text{Flip}_c = \emptyset].$$

Proof. Identical to Theorem 4.2; only the bookkeeping changes. Let $\pi' = (x_{c+i} y_i) \circ \pi$. The preimage chain of 1 under π' agrees with that of π up to depth $c + i$ and then continues along the chain of 2: $\pi'^{-(c+i+t)}(1) = y_{i+t}$; symmetrically $\pi'^{-(i+t)}(2) = x_{c+i+t}$ ($t \geq 1$). So the pair of D_c at y -depth $i' > i$ is $\{y_{i'}, x_{c+i'}\}$, unchanged as a set, and the pairs at depth $i' < i$ are untouched: D_c is preserved. (The same computation shows that trading a pair at depths (k, k') sends a deeper pair at depths (K, K') to depths $(K' + k - k', K - k + k')$, whose offset is $2(k - k') - (K - K')$; it is preserved iff $K - K' = k - k'$, so only families of constant offset, i.e. the D_c , are preserved.) On A the trade merges C_1, C_2 into a cycle with $d'_{12} = |C_1| - (c + i) + i = |C_1| - c$ and $d'_{21} = |C_2| + c$, so $K'_c = \min(d'_{21} - c, d'_{12}) = \min(|C_2|, |C_1| - c) = K_c$; on B the pair separates rows 1, 2 ($c + i < d_{21}$ puts x_{c+i} on the arc from 2 to 1, $i < d_{12}$ puts y_i on the arc from 1 to 2) and gives $|C'_1| = d_{12} + c$, $|C'_2| = d_{21} - c$, again $K'_c = \min(d_{12}, d_{21} - c) = K_c$. No pair of D_c is frame-adjacent: $y_i = \pi(x_{c+i})$ would need $d_{21} = c - 1$ (or $\pi(1) = 2$ when $c = 0$) and $x_{c+i} = \pi(y_i)$ would need $d_{21} = c + 1$, both excluded on $\mathcal{D}_c$. Commutation, preservation of flippability and the orbit count are as before. $\square$

Exact enumeration at $n = 7$ ($\lambda = (7)$, $\{\alpha, \beta\} = \{2, 5\}$) confirms the identities for $c = 0, 1, 2, 3$ to all digits ($\mathbb{P}[B, \mathcal{D}_c] - \mathbb{P}[A, \mathcal{D}_c] = -0.005848, +0.001462, +0.003509, -0.001316$, equal to the right-hand sides).

Using the offsets to price short arcs. Fix $\ell \geq 2$ and consider B -instances with $d_{12} = \ell$. For $0 \leq c \leq n/2 - \ell$ and $d_{21} \geq c + \ell$ one has $K_c = \ell$, so D_c has $\ell - 1$ pairs, and the offset- c trade at a flippable pair of D_c lands in A with $|C_1| = \ell + c \leq n/2$ and $|C_2| = d_{21} - c \geq \ell$. Within a class of $\mathcal{D}_c$ all B -members have the same (d_{12}, d_{21}) and all A -members the same $(|C_1|, |C_2|) = (d_{12} + c, d_{21} - c)$, and Flip_c is a class invariant, so Theorem 4.5 class by class and Proposition 4.6 give

$$\mathbb{P}[B, d_{12} = \ell, d_{21} \geq c + \ell, \text{Flip}_c \neq \emptyset] = \mathbb{P}[A, |C_1| = \ell + c, |C_2| \geq \ell, \text{Flip}_c \neq \emptyset] \leq \frac{18}{n - \ell - c - 1} \leq \frac{36}{n - 2}. \quad (3)$$

Averaging over c : if $\theta(L, P)$ denotes the fraction of offsets $c \in [0, n/2 - \ell]$ whose diagonal D_c contains a flippable pair, then for every $\theta_0 > 0$

$$\mathbb{P}[B, d_{12} = \ell, |C_1| \geq n/2 + \ell, \theta \geq \theta_0] \leq \frac{36}{(n - 2)\theta_0}, \quad (4)$$

since on $|C_1| \geq n/2 + \ell$ one has $d_{21} = |C_1| - \ell \geq n/2 \geq c + \ell$ for every such c , so $\mathbb{E}[\theta \mathbf{1}_E]$ is the average over c of the left side of (3), and Markov's inequality applies. Together with the splice-in bound $\mathbb{P}[B, d_{12} = \ell, |C_1| < n/2 + \ell] \leq 4/(n - 2\ell + 1) \leq 8/n$ for $\ell \leq n/4$ (Proposition 4.7 below), this prices every short arc except on the event $\{\theta < \theta_0\}$.

Proposition 4.6 (short frame cycles through a marked row are rare). *For (L, P) uniform on X and every $2 \leq \ell \leq n/2$,*

$$\mathbb{P}[|C_1| = \ell] \leq \frac{18}{n - \ell - 1}, \quad \text{hence} \quad \mathbb{P}[|C_1| \leq \ell] \leq \frac{36\ell}{n} \quad (\ell \leq n/2 - 1),$$

and the same bounds hold for C_2 .

Proof. Let $E_\ell = \{(L, P) \in X : |C_1| = \ell\}$. (The case $C_1 = \{1, 2\}$ cannot occur: it would mean $\sigma(j) = j'$ and $\sigma(j') = j$, but $j' = \sigma^\alpha(j)$ lies in an m -cycle with $m \geq 4$.)

Forward moves. From $(L, P) \in E_\ell$ choose $y \in C_1 \setminus \{1, 2\}$ and $x \notin C_1 \cup \{2\}$: at least $(\ell - 2)(n - \ell - 1)$ pairs $((\ell - 1)(n - \ell - 1)$ if $2 \notin C_1$; for $\ell = 2$ this is $n - 3$). If (x, y) is flippable, apply the j' -trade at (x, y) (trade rows x, y along the $\rho_{x,y}$ -cycle $Q_{x,y}$ through j'); otherwise turn C_1 (which makes (x, y) flippable, as $y \in C_1 \not\equiv x$) and then trade $Q_{x,y}$ in the turned square. In both cases the resulting frame is $(xy) \circ \pi^\pm$ with $\pi^\pm$ equal to π or to π inverted on C_1 , so the cycle of row 1 becomes $C_1 \cup C_x$, of length $\ell + |C_x| > \ell$; rows 1, 2 are untouched by the trade, and the result is in X (no turn, or $2 \in C_1$) or in J (turn with $2 \notin C_1$).

Backward moves. A square $(L', P) \in X \cup J$ is reached from $(L, P) \in E_\ell$ by the move at (x, y) only if x, y lie in the cycle C'_1 of row 1 in L' , and undoing the trade (the same trade, since $\rho'_{x,y} = \rho_{x,y}^{-1}$ on $Q_{x,y}$) splits C'_1 into the cycles $\{x, \pi'(x), \dots, \pi'^{-1}(y)\}$ and $\{y, \pi'(y), \dots, \pi'^{-1}(x)\}$ of $(xy) \circ \pi'$, of which the one containing row 1 must have length exactly ℓ : given L' , that arc is determined by its position relative to row 1 (ℓ choices), and the arc determines $\{x, y\}$. With the binary choice “turn back or not” this gives at most 2ℓ preimages per (L', P) .

Counting. The map $((L, P), \text{move}) \mapsto ((L', P), \text{inverse move})$ is injective, so $|E_\ell|(\ell - 2)(n - \ell - 1) \leq 2\ell(|X| + |J|) \leq 6\ell|X|$, using $|J| = 2|X_A| \leq 2|X|$ (Theorem 2.4). For $\ell \geq 3$ this is the

claim; for $\ell = 2$ the forward count is $n - 3$ and the backward count 4, giving $12/(n - 3)$. The bound for C_2 is proved by the same argument with the roles of rows 1 and 2 exchanged: a turn of $C_2 \not\equiv 1$ replaces σ by $(\sigma(j) \sigma(j')) \circ \sigma$, which also splits the m -cycle into lengths α, β and lands in J . $\square$

Proposition 4.7 (the splice-in bound). *For $2 \leq \ell \leq n/4$, $\mathbb{P}_X[B, d_{12} = \ell, |C_1| < n/2 + \ell] \leq 4/(n - 2\ell + 1) \leq 8/n$, and the same with d_{21} in place of d_{12} .*

Proof. Let E be the event. *Forward moves.* From $(L, P) \in E$ choose $y = \pi^i(1)$ with $1 \leq i \leq \ell - 1$ (a row strictly inside the arc from 1 to 2; $\ell - 1$ choices) and a row $x \notin C_1$ ($n - |C_1| \geq (n - 2\ell + 1)/2$ choices), and let C_x be the frame cycle of x , which contains neither 1 nor 2. If (x, y) is not flippable, turn C_x : this trades columns j, j' on the rows of C_x only, so it preserves X , inverts π on C_x (Lemma 2.2), fixes row y , and replaces $\rho_{x,y}$ by $\rho_{x,y} \circ (j j')$, which makes (x, y) flippable. Then apply the j' -trade at (x, y) , which replaces the frame by $(x y) \circ \pi^\pm$ ($\pi^\pm = \pi$, or π inverted on C_x): the arc from 1 to 2 becomes $1 \rightarrow \dots \rightarrow \pi^{-1}(y) \rightarrow x \rightarrow \dots \rightarrow y \rightarrow \dots \rightarrow 2$ with C_x occupying positions $i, \dots, i + |C_x| - 1$, so d_{12} grows by $|C_x|$, rows 1, 2 are untouched, and the result is again in X_B with the same mark. *Backward moves.* Given the target $(L', P) \in X_B$ and ℓ , a preimage is determined by the position $i \in \{1, \dots, \ell - 1\}$ of x on the arc of L' from 1 to 2 (then $|C_x| = d'_{12} - \ell$ and $y = \pi'^{i+|C_x|}(1)$ are determined), together with the binary choice of whether C_x was turned; undoing is the same j' -trade (the pair stays flippable, $\rho'_{x,y} = \rho_{x,y}^{-1}$ on the traded cycle) followed by the turn. So each target has at most $2(\ell - 1)$ preimages while each source has at least $(\ell - 1)(n - 2\ell + 1)/2$ forward moves, whence $|E|(\ell - 1)(n - 2\ell + 1)/2 \leq 2(\ell - 1)|X|$. $\square$

5 The witness hypothesis

Notation as in §4: $(L, P) \in X$, $P = \{j, j'\}$, frame π , ladder pairs (x_k, y_k) , K_0 , and for B -instances the offset diagonals D_c of §4.1. The statement below is the reference hypothesis of this paper — the least demanding sufficient condition for (L) that we have; the orbit method of §6 gives two others.

Say that a pair of rows (x, y) , $x, y \notin \{1, 2\}$, carries a *witness* if the $\rho_{x,y}$ -cycle through column j has length at most $\ell' = \log^5 n$ and does not contain j' (then (x, y) is flippable). A frame-adjacent pair ($y = \pi(x)$, so $\rho_{x,y}(j) = j'$) never carries one; the families below contain no such pair. For a pair of rows of a random Latin square the length of the j -cycle is spread like a uniform variable on $\{2, \dots, n\}$ (Remark 5.4 below), so a witness has probability $\approx \ell'/n$.

Hypothesis 5.1 (witness lemma). *Let (L, P) be uniform on X , uniformly in λ, α, β . Let $\mathcal{G}$ be one of the following families of blocks of pairs of rows, determined by the frame: (i) the ladder, each pair a block of its own; (ii) for a B -instance, with $\ell := d_{12}$, provided $\ell \leq n/4$ and $|C_1| \geq n/2 + \ell$, the diagonals D_c , $0 \leq c \leq n/2 - \ell$, each a block of $\ell - 1$ pairs (one family, not one per ℓ). Let $G_{\mathcal{G}}$ be the number of blocks containing a pair that carries a witness, and $\mu_{\mathcal{G}} = \sum_{B \in \mathcal{G}} \min(\frac{1}{2}, \ell'|\mathcal{B}|/n)$. Then, for the ladder (i),*

$$\mathbb{P}[G_{\mathcal{G}} = 0 \text{ and } \mu_{\mathcal{G}} \geq \log^2 n] = o(1/\log n),$$

and for the diagonals (ii) there is a constant $\kappa_1 > 0$ such that

$$\mathbb{P}[G_{\mathcal{G}} < \kappa_1 \mu_{\mathcal{G}} \text{ and } \mu_{\mathcal{G}} \geq \log^2 n] = o(1/\log n).$$

(This is annealed: an exceptional set of frames of probability $o(1/\log n)$ is allowed. For a single pair of non-adjacent rows the witness probability is $\approx \ell'/n$, Remark 5.4, so $\kappa_1 \mu_G$ is a constant fraction of the expected count if blocks are nearly independent given their shared rows.)

Proposition 5.2 ((L) follows from the witness lemma). *Hypothesis 5.1 implies (L), hence the weak CGW conjecture.*

Proof. Put $\ell_0 = n/\log^2 n$. By Theorem 4.2, $|\mathbb{P}[B] - \mathbb{P}[A]| \leq \mathbb{P}[\text{Flip} = \emptyset] \leq \mathbb{P}[K_0 \geq \ell_0, \text{Flip} = \emptyset] + \mathbb{P}[A, K_0 < \ell_0] + \mathbb{P}[B, K_0 < \ell_0]$.

Long ladders. On $\{K_0 \geq \ell_0\}$ take family (i): $\mu \geq (\ell_0 - 1)\ell'/n \geq \log^2 n$, so $G \geq 1$ except on an event of probability $o(1/\log n)$, and $G \geq 1$ means $\text{Flip} \neq \emptyset$.

Short cycles. $\mathbb{P}[A, K_0 < \ell_0] \leq 72\ell_0/n = 72/\log^2 n$ by Proposition 4.6.

Short arcs. $\mathbb{P}[B, K_0 < \ell_0] \leq \sum_{\ell < \ell_0} (\mathbb{P}[B, d_{12} = \ell] + \mathbb{P}[B, d_{21} = \ell])$; we treat $d_{12} = \ell$; the case $d_{21} = \ell$ reduces to it by exchanging rows 1 and 2 (with the symbol relabelling that restores the first row): this replaces σ by σ^{-1} , keeps the mark $\{j, j'\}$ with $j' = \sigma^{-\beta}(j)$, so maps $X(n, \lambda; \alpha, \beta)$ onto $X(n, \lambda; \beta, \alpha)$ measure-preservingly, conjugates the frame by (1 2), hence exchanges d_{12} with d_{21} , preserves the ladder as a set of unordered pairs and carries the diagonals with the x -row deeper to those with the y -row deeper, and leaves every $\rho_{x,y}$ with $x, y \notin \{1, 2\}$ unchanged; so the $d_{21} = \ell$ case for (λ, α, β) is the $d_{12} = \ell$ case of the hypothesis for (λ, β, α) , which “uniformly in λ, α, β ” covers (for $\alpha = \beta$ the hypothesis is read for the ordered mark). Fix ℓ . The part with $|C_1| < n/2 + \ell$ costs $\leq 8/n$ (Proposition 4.7). On $|C_1| \geq n/2 + \ell$ take family (ii), with $N = \lfloor n/2 - \ell \rfloor + 1 \geq n/4$ blocks: $\mu = N \min(\frac{1}{2}, (\ell - 1)\ell'/n) \geq \frac{1}{4} \min(n/2, (\ell - 1)\ell') \geq \log^2 n$, so except on an event F_ℓ of probability $o(1/\log n)$ the fraction of offsets whose diagonal contains a flippable pair is $\theta = G/N \geq \theta_0(\ell) := \kappa_1 \min(\frac{1}{2}, (\ell - 1)\ell'/n)$. By (4),

$$\begin{aligned} \sum_{\ell < \ell_0} \mathbb{P}[B, d_{12} = \ell, |C_1| \geq n/2 + \ell, \theta \geq \theta_0(\ell)] &\leq \sum_{\ell - 1 \leq n/(2\ell')} \frac{36n}{(n-2)\kappa_1(\ell-1)\ell'} + \sum_{n/(2\ell') < \ell - 1 < \ell_0} \frac{72}{(n-2)\kappa_1} \\ &= O\left(\frac{\log n}{\ell'} + \frac{\ell_0}{n}\right) = O\left(\frac{1}{\log^2 n}\right), \end{aligned}$$

and $\sum_{\ell < \ell_0} 8/n \leq 8/\log^2 n$. The exceptional events: the family (ii) is a function of the frame, and the events $\{d_{12} = \ell\}$ are disjoint, so $\sum_{\ell} \mathbb{P}[B, d_{12} = \ell, |C_1| \geq n/2 + \ell, F_\ell] = \mathbb{P}[\bigcup_{\ell} \{B, d_{12} = \ell, |C_1| \geq n/2 + \ell, F_\ell\}]$ is the probability of a single failure event of the hypothesis, which is $o(1/\log n)$. Altogether $\mathbb{P}[\text{Flip} = \emptyset] = o(1/\log n)$. $\square$

Remark 5.3 (what the witness lemma is, and is not). (i) It is a statement about substructures of polylogarithmic size (a cycle of length $\leq \log^5 n$ of a row-pair permutation through a given column), in the space of completions of two rows and two columns, for families of pairs selected by the frame. It asks for a lower bound on their expected number and a lower-tail concentration (a second moment with relative variance $o(1/\log n)$ suffices). For two *fixed* rows, counts of such substructures are what the switching method of [10] controls; what it does not address is a family of row pairs selected by the square, and that is where the difficulty sits (§9). (ii) It asks nothing about long cycles, same-cycle events, or the fairness of any global statistic. (iii) Rigid squares ($\mathbb{Z}_p$ isotopes and their relatives, where no row pair has a short cycle) enter only through the failure probability, which must be $o(1/\log n)$; for them the hypothesis is simply the statement that they are rare. (iv) The constants: $\ell' = \log^5 n$ and $\ell_0 = n/\log^2 n$ are one admissible choice; the proof needs $\ell_0 \ell'/n \geq \log^2 n$, $\ell' \geq 4 \log^2 n$, $\ell_0/n = o(1/\log n)$ and $\log n/\ell' = o(1/\log n)$. (v) The

block form is what the short-arc argument uses: the number of offsets with a witness, not the number of witnesses, since the diagonals of different offsets are dependent through the shared rows y_i . (vi) The hypothesis is stated only for the two families used; a version for arbitrary frame-determined families would be false, since frame-adjacent pairs never carry a witness.

Remark 5.4 (data). On Jacobson–Matthews samples with a uniformly random mark (`runs/s13/f1/`), the length of the j -cycle of ρ_{x_k, y_k} for a ladder pair is spread like a uniform variable on $\{2, \dots, n\}$: $\mathbb{P}[\leq \ell] = 0.021, 0.081, 0.18, 0.38$ for $\ell = 2, 5, 10, 20$ at $n = 50$ against $(\ell - 1)/(n - 1) = 0.020, 0.082, 0.18, 0.39$, and $0.010, 0.041, 0.092, 0.19$ at $n = 100$; among instances with $K_0 \geq n/4$ a flippable ladder pair with j -cycle of length ≤ 10 exists in 93% ($n = 50$) and 95% ($n = 100$) of the instances, and one of length ≤ 20 in 99.6% and 99.8%. These samples see only typical types λ ; the hypothesis is required uniformly in λ . The fixed-rectangle sampler of Remark 6.14 reaches the rare types: there the witness events along the ladder are independent Bernoulli variables to three digits (flat hazard, $\mathbb{P}[\text{no witness in } K \text{ pairs}] = (1 - w_1)^K$ to $K = 40$ at $n = 100$) and the witness counts across the offset diagonals have binomial lower tails, identically for (n) and for $(4, 2^{(n-4)/2})$. Classes with all cycles of $\rho_{1,2}$ of length ≤ 3 carry no admissible mark and do not enter.

6 The orbit method: alternative sufficient conditions

Throughout this section $X = X(n, \lambda; \alpha, \beta)$, and “legal” means: preserves the type of rows 1, 2 and the mark, i.e. maps X to itself. The mark is now written $\{p, p'\}$ (so $p = j$, $p' = j'$ in the notation of §§2–4; the letters j, j' are freed for other columns), π is its frame, and a column pair (q, c) is *free* if $\{q, c\} \cap \{p, p'\} = \emptyset$. A pair of rows (x, y) is called *parallel* if it is flippable (p, p' in different cycles of $\rho_{x,y}$) and *crossed* otherwise; Q_p and $Q_{p'}$ denote the $\rho_{x,y}$ -cycles through p and p' . A free pair (q, c) is *separated* by $\{p, p'\}$ in $\rho_{x,y}$ if q and c lie on different arcs of a common cycle through p, p' (crossed case) or one in Q_p and the other in $Q_{p'} \neq Q_p$ (parallel case). The method bounds $\mathbb{P}_X[\text{Flip}_I = \emptyset]$ by a product of conditionally independent factors; the two hypotheses it leads to are compared with Hypothesis 5.1 after (c'), and in (e) the same mechanism, applied to the first ladder pair alone, gives an exact expression for $\mathbb{P}_X[B] - \mathbb{P}_X[A]$ and a third alternative hypothesis.

(a) Transposition: which half of the conditioning is the obstacle. Let L be uniform on *all* Latin squares of order n (no condition on rows 1, 2 in this paragraph), let x, y carry four distinct symbols on $\{p, p'\}$, and condition on the two full columns p, p' (the frame π of (p, p')). In the transposed square $L^\top$ the columns p, p' are the distinguished rows, the frame π is their column permutation, and the event “ p, p' lie in different cycles of $\rho_{x,y}$ ” is the event “ $\{x, y\}$ is an A -pair for the distinguished rows”. Conditioning on the two full rows p, p' of $L^\top$ is conditioning on their type and on the position of the pair $\{x, y\}$ in their cycle structure, which is free by symbol and column relabelling. Hence, by Lemma 2.3 ($|J_A| = |J_B|$, which needs no condition on the cycle lengths) and CGW’s Theorem 3.13 (the inequality $\mathbb{P}_X[A] \geq \frac{1}{3}$, whose published proof has a gap for every split other than $(2, 2)$, Remark 3.5):

$$\mathbb{P}[p, p' \text{ in different cycles of } \rho_{x,y} \mid \text{columns } p, p'] \begin{cases} = \frac{1}{2} & \text{if } x, y \text{ are apart in the frame,} \\ \geq \frac{1}{3} & \text{if together at distances } \geq 2. \end{cases} \quad (5)$$

(The second bound is CGW’s $\frac{3}{2}$, affected by the gap; it holds with $\frac{1}{4}$ unconditionally at distance 2, Remark 3.5.) So the statement with the *columns* prescribed is CGW’s theorem; it is the

conditioning on rows 1, 2 — on their type and the mark, which is what defines X and cannot be removed — that is the obstacle. In the same transposed picture the ladder trades T_k of Theorem 4.2 at flippable pairs are CGW's flips and backflips at the column pairs $\{x_k, y_k\}$ of $L^\top$, and Theorem 4.2 is exactly the legal half of CGW's Lemma 3.12 for the distinguished rows p, p' : the other half — the switch and the cross-switch at $\{x_k, y_k\}$ — trades the columns p, p' of L along a frame cycle or frame arc through exactly one of rows 1, 2 (apart: $C_1 \ni 1$, $C_1 \not\ni 2$; together: the arc from x_k to y_k contains 1 but not 2, the other arc 2 but not 1), so it changes the type of rows 1, 2 and leaves X . The switch is, in the original orientation, the unflip of rows 1, 2 at (p, p') (which maps $X_A \rightarrow J_A$); it toggles the crossed/parallel status of *every* pair (x, y) with x on the side of row 1 and y on the side of row 2, and reverses the orientation of one side, so it carries the ladder (a diagonal of the side-by-side array of pairs) to a different family (an anti-diagonal). The cross-switch likewise maps X_B to J_B (it splits the σ -cycle through p, p' while keeping rows 1, 2 in one frame cycle, CGW Lemma 3.10) and toggles the interior pairs. Both are exact but not usable for the ladder bits in X .

(b) Row trades are free in X ; column turns are conditionally free. For rows $z, w \notin \{1, 2\}$ every trade of rows z, w along a $\rho_{z,w}$ -cycle or $\rho_{z,w}$ -arc changes rows z, w only, hence is legal in X . In particular:

Lemma 6.1 (exact fairness of column-pair status for generic rows). *Let $x, y \notin \{1, 2\}$ and let (q, c) be a column pair lying in two different cycles of $\rho_{x,y}$. The CGW switch at (q, c) for the rows x, y (trade of rows x, y along the one of the two cycles that contains the smaller minimum symbol) is a legal involution of X that toggles whether x, y lie in one column cycle of (q, c) , and preserves the cycle partition of $\rho_{x,y}$ and the separation status of every column pair with respect to $\{p, p'\}$. Hence, conditionally on everything determined by the other rows and by the cycle partition of $\rho_{x,y}$,*

$$\mathbb{P}_X[x, y \text{ in different column cycles of } (q, c)] = \frac{1}{2}.$$

Proof. As in Lemma 2.3 with $(1, 2)$ replaced by (x, y) : the trade swaps the entries of rows x, y on the columns of the chosen cycle $Q \ni q$, $Q \not\ni c$, so $\delta_{q,c} \mapsto \delta_{q,c} \circ (xy)$, which splits or joins the column cycles through x, y . It inverts Q in place, so the symbol sets of the two cycles are unchanged and the min-symbol rule re-selects Q : an involution. Rows 1, 2 are untouched, so type and mark are preserved. Inverting Q exchanges the two arcs between p and p' when $Q \ni p, p'$, so separation of any pair by $\{p, p'\}$ is preserved. (If Q contains p or p' the frame changes at rows x, y ; so frame-determined data are part of the conditioning only when Q avoids p, p' .) $\square$

A column turn, by contrast, changes the rows of a column cycle. Let (q, c) be a free column pair and Z one of its column cycles.

Lemma 6.2 (legal column turns). *The turn of Z (swapping the entries of columns q, c in the rows of Z) is a measure-preserving involution of X if and only if $|Z \cap \{1, 2\}| \in \{0, 2\}$; when $|Z \cap \{1, 2\}| = 2$ it conjugates σ by (qc) and keeps the type and the mark. It preserves the columns p, p' (hence the frame and every frame-determined family of row pairs) and the column-cycle partition of (q, c) and of every column pair disjoint from $\{q, c\}$ (in particular of every pair of a matching containing (q, c)). For a row pair (x, y) with $|Z \cap \{x, y\}| = 1$ it replaces $\rho_{x,y}$ by $\rho_{x,y} \circ (qc)$ (if $x \in Z$) or $(qc) \circ \rho_{x,y}$ (if $y \in Z$), which toggles the crossed/parallel status of (x, y) exactly when (q, c) is separated by $\{p, p'\}$ in $\rho_{x,y}$ (different arcs when crossed, different*

cycles $Q_p, Q_{p'}$ when parallel), and leaves (q, c) separated afterwards; with $|Z \cap \{x, y\}| \in \{0, 2\}$ the status is unchanged. Turns of cycles of column pairs from a fixed matching $\mathcal{M}$ of the free columns pairwise commute.

Proof. A column cycle is a Latin trade (Lemma 2.2), so the turn is a bijection of Latin squares and an involution, and it preserves the cycle partition of (q, c) ; column pairs disjoint from $\{q, c\}$ are untouched, and so are columns p, p' . Rows 1, 2 change iff Z meets them; if $Z \ni 1, 2$ both rows swap their entries at q, c , so $\sigma' = (qc)\sigma(qc)$, whose type is that of σ and in which p, p' (fixed by (qc)) keep their distance. If Z meets exactly one of rows 1, 2, $\sigma' = \sigma \circ (qc)$ or $(qc) \circ \sigma$, which changes the type. The effect on $\rho_{x,y}$: if $x \in Z$ and $y \notin Z$ the entries of row x at q, c are exchanged, so $\rho'_{x,y} = \rho_{x,y} \circ (qc)$, and symmetrically; multiplying by a transposition of two points on one cycle splits it at those points, so p, p' are separated iff q, c lie on different arcs between them, and multiplying by a transposition of two points on different cycles merges the cycles, so p, p' are joined iff q, c lie one in each of their cycles; in both cases q, c lie one on each side afterwards. Turns of cycles of disjoint column pairs act on disjoint sets of cells; two cycles of the same pair have disjoint row sets. $\square$

(c) The product bound. Fix a matching $\mathcal{M}$ of the free columns, independent of the square, and a frame-determined family $\mathcal{G}$ of row pairs (x, y) , $x, y \notin \{1, 2\}$, whose x -rows are pairwise distinct and disjoint from its y -rows (a block of the ladder, where the pairs are even disjoint; or a union of diagonals D_c with distinct x -rows, which may share y -rows). For $(x, y) \in \mathcal{G}$ call a column pair $(q, c) \in \mathcal{M}$ a *clean coordinate* of (x, y) if the column cycle Z of (q, c) through x has length at most ℓ_1 (a fixed constant), satisfies $|Z \cap \{1, 2\}| \in \{0, 2\}$, and contains no row of $\mathcal{G}$ other than x (no other x -row and no y -row, in particular not y). Let $\mathcal{T}_{x,y}$ be the set of clean coordinates of (x, y) and $\eta_{x,y} \in \{0, 1\}^{\mathcal{T}_{x,y}}$. By Lemma 6.2 the turns of all clean coordinates of all pairs of $\mathcal{G}$ commute, generate a group acting freely on X , preserve $\mathcal{G}$ and the sets $\mathcal{T}_{x,y}$, and

$$\rho_{x,y}(\eta) = \rho_{x,y}(0) \circ \prod_{(q,c) \in \mathcal{T}_{x,y}: \eta_{x,y}(q,c)=1} (qc),$$

the bit $b_{x,y}(\eta) = [(x, y) \text{ parallel}]$ depending on $\eta_{x,y}$ only. Let the *orbit weight* of (x, y) be

$$\bar{y}_{x,y} = \mathbb{E}_{\eta_{x,y}} \left[\frac{\#\{(q, c) \in \mathcal{T}_{x,y} : (q, c) \text{ separated by } \{p, p'\} \text{ in } \rho_{x,y}(\eta)\}}{|\mathcal{T}_{x,y}|} \right]$$

(0 if $\mathcal{T}_{x,y} = \emptyset$), a function of the orbit, and $Y_{\mathcal{G}} = \sum_{(x,y) \in \mathcal{G}} \bar{y}_{x,y}$.

Proposition 6.3 (product bound). *Conditionally on the orbit, the bits $b_{x,y}$, $(x, y) \in \mathcal{G}$, are independent, and*

$$\mathbb{P}[(x, y) \text{ crossed} \mid \text{orbit}] \leq 1 - \frac{1}{2} \bar{y}_{x,y}, \quad \mathbb{P}[\text{no pair of } \mathcal{G} \text{ is flippable} \mid \text{orbit}] \leq \prod_{(x,y) \in \mathcal{G}} (1 - \frac{1}{2} \bar{y}_{x,y}) \leq e^{-Y_{\mathcal{G}}/2}.$$

Consequently $\mathbb{P}_X[\text{Flip}_{\mathcal{G}} = \emptyset] \leq \mathbb{E}_X[e^{-Y_{\mathcal{G}}/2}]$.

Proof. The action is free: distinct column pairs act on disjoint cells, and two clean coordinates of different pairs at the same (q, c) are the cycles through two distinct x -rows, each avoiding the other, hence different cycles; the sets $\mathcal{T}_{x,y}$ are orbit-invariant by Lemma 6.2 (row sets of the cycles, their meeting $\{1, 2\}$, and the frame-determined family are preserved). The uniform

measure on X restricted to an orbit is therefore uniform on the orbit, which is the product of the cubes $\{0,1\}^{\mathcal{T}_{x,y}}$. A clean coordinate of (x,y) contains x and not y , and a clean coordinate of another pair contains neither x nor y (it avoids all x -rows but its own and all y -rows), so $b_{x,y}$ depends on $\eta_{x,y}$ alone, whence the independence. Fix (x,y) and $Z = (q,c) \in \mathcal{T}_{x,y}$. Whether Z is separated in $\rho_{x,y}(\eta)$ does not depend on η_Z (Lemma 6.2: the turn of Z leaves Z separated), and when it is, the two values of η_Z give opposite bits. Hence $\mathbb{P}[b_{x,y} = 0] \leq 1 - \frac{1}{2}\mathbb{P}_\eta[Z \text{ separated}]$ for every Z , and averaging the right side over $Z \in \mathcal{T}_{x,y}$ gives the first bound; the second is the product of the first over the independent bits. $\square$

Hypothesis 6.4 (orbit genericity; fixed matching). *Fix the matching $\mathcal{M}$ of the free columns in consecutive pairs (in the natural order of the columns) and a constant ℓ_1 as in (c), and for a family $\mathcal{G}$ and $D \subseteq \mathcal{G}$ write $Y_D = \sum_{(x,y) \in D} \bar{y}_{x,y}$ (orbit weights computed with cleanliness with respect to all rows of $\mathcal{G}$). There is a constant $\kappa_0 > 0$ such that, with $\ell_0 = n/\log^2 n$ and uniformly in λ, α, β : (i) for the sub-ladder I of the first $\min(K_0 - 1, \lceil n/8 \rceil)$ ladder pairs, $\mathbb{E}_X[e^{-Y_I/2}; K_0 \geq \ell_0] = o(1/\log n)$; (ii) for a B -instance, with $\ell := d_{12} < \ell_0$ and $|C_1| \geq n/2 + \ell$, $S = \{2\ell t : 0 \leq t \leq (n/2 - \ell)/(2\ell)\}$ and $\mathcal{G} = \bigcup_{c \in S} D_c$: the event that fewer than $\kappa_0|S|$ of the diagonals D_c , $c \in S$, have $Y_{D_c} \geq \kappa_0(\ell - 1)$ has probability $o(1/\log n)$ (one event, not one per ℓ); and symmetrically for d_{21} . (A fixed small fraction, not a half: with a fixed matching the orbit weight of a single pair vanishes with probability ≈ 0.7 , §(d), so for $\ell = 2$ only a minority of the one-pair diagonals can qualify.)*

Since $\bar{y}_{x,y}$ is typically of constant size (§(d): 0.16–0.19 on average with the fixed matching), (i) asks that the sum of $\approx n/8$ such weights not collapse; its content is a *genericity constant* — that a clean coordinate is separated by the mark with probability bounded below — which none of the exact tools of this paper provides (§9).

(c') Adaptive coordinates: the toggling is automatic. Proposition 6.3 needs the clean coordinates of a ladder pair to be *separated* by $\{p, p'\}$ with positive probability, a genericity input we cannot prove (§9). It disappears if the coordinates are chosen along $\rho_{x,y}$ from the mark. For a pair (x,y) of the family and $t \geq 1$ let

$$(q^t, c^t) := (\rho_{x,y}^t(p), \rho_{x,y}^t(p')),$$

and let $\mu_{x,y} := \min(d(p \rightarrow p'), d(p' \rightarrow p))$ if (x,y) is crossed and $\min(|Q_p|, |Q_{p'}|)$ if parallel. For $t < \mu_{x,y}$ the pair (q^t, c^t) is free ($q^t \neq c^t$, both $\notin \{p, p'\}$) and separated by $\{p, p'\}$ in $\rho_{x,y}$ by construction: on different arcs when crossed, one in Q_p and one in $Q_{p'}$ when parallel. So the turn of its column cycle Z through x (with $y \notin Z$) toggles the status of (x,y) whenever it is legal (Lemma 6.2). What makes these coordinates usable in the orbit method is that they are invariant under their own turn. (We keep (c) for comparison: its orbit weight is what the data of (d) measure, and it is where the genericity constant of §9 lives.)

Lemma 6.5 (the adaptive coordinate is orbit-invariant). *Let $t < \mu_{x,y}$, let Z be the column cycle of (q^t, c^t) through x with $y \notin Z$, and let L' be the turn of Z . Then $\rho'_{x,y} = \rho_{x,y} \circ (q^t c^t)$, the status of (x,y) toggles, $\mu'_{x,y} = \mu_{x,y}$, and $(\rho'^s(p), \rho'^s(p')) = (q^s, c^s)$ for every $s \leq t$; in particular the t -th candidate of L' is (q^t, c^t) again, and its column cycle through x is Z (same row set).*

Proof. Write $\rho = \rho_{x,y}$, $q = q^t$, $c = c^t$. The turn exchanges the entries of x at q, c (and those of the other rows of Z , which do not affect ρ), so $\rho' = \rho \circ (q c)$: $\rho'(q) = \rho(c)$, $\rho'(c) = \rho(q)$, $\rho' = \rho$

elsewhere. Crossed case, with $p = u_0 \rightarrow u_1 \rightarrow \dots \rightarrow u_a = p' \rightarrow v_1 \rightarrow \dots \rightarrow v_b = p$ and $q = u_t$, $c = v_t$ ($t < \min(a, b)$): the cycles of ρ' through p and p' are $u_0 \rightarrow \dots \rightarrow u_t \rightarrow v_{t+1} \rightarrow \dots \rightarrow v_b = p$ and $u_{t+1} \rightarrow \dots \rightarrow u_a = p' \rightarrow v_1 \rightarrow \dots \rightarrow v_t \rightarrow u_{t+1}$, of lengths b and a , so (x, y) is now parallel with $\mu' = \min(a, b) = \mu$, and the paths $p \rightarrow u_s, p' \rightarrow v_s$ for $s \leq t$ are untouched. Parallel case, $Q_p = (p = u_0, u_1, \dots, u_{a-1})$, $Q_{p'} = (p' = v_0, \dots, v_{b-1})$, $q = u_t$, $c = v_t$ ($t < \min(a, b)$): the single cycle of ρ' is $u_0 \rightarrow \dots \rightarrow u_t \rightarrow v_{t+1} \rightarrow \dots \rightarrow v_{b-1} \rightarrow v_0 = p' \rightarrow v_1 \rightarrow \dots \rightarrow v_t \rightarrow u_{t+1} \rightarrow \dots \rightarrow u_{a-1} \rightarrow u_0$, with $d'(p \rightarrow p') = b$, $d'(p' \rightarrow p) = a$, so crossed with $\mu' = \mu$, and again the paths up to step t are untouched. The row set of the turned cycle is preserved (Lemma 2.2). $\square$

Proposition 6.6 (orbit bound with adaptive coordinates). *Let $(L, P) \in X$, I a frame-determined family of pairs (x_k, y_k) as in §(c) (distinct x -rows, disjoint from the y -rows, none in $\{1, 2\}$; the ladder, or a union of offset diagonals), and $T, \ell_1 \geq 1$. Call the candidate (k, t) good if $t \leq \min(T, \mu_k - 1)$ and the column cycle Z_k^t of (q_k^t, c_k^t) through x_k has length $\leq \ell_1$, meets $\{1, 2\}$ in 0 or 2 rows and meets the rows of I only in x_k . Let $t_0(k)$ be the least good t (if any), $C_k = \{q_k^t, c_k^t : t \leq t_0(k)\}$ (all candidates $t \leq \min(T, \mu_k - 1)$ if there is none), and*

$$U = \{k \in I : t_0(k) \text{ exists and } \{q_k^{t_0(k)}, c_k^{t_0(k)}\} \cap C_l = \emptyset \text{ for all } l \in I \setminus \{k\}\}.$$

Then the turns T_k of $Z_k^{t_0(k)}$, $k \in U$, are legal, commute, generate a free action of $\{0, 1\}^U$ on X that preserves the frame, the family, every μ_l , $t_0(l)$, C_l ($l \in I$) and U ; T_k toggles the status of (x_k, y_k) and fixes that of every other pair. Hence

$$\mathbb{P}_X[\text{Flip}_I = \emptyset] \leq \mathbb{E}_X[2^{-|U|}].$$

Proof. T_k changes only columns $q_k, c_k := q_k^{t_0(k)}, c_k^{t_0(k)}$ in the rows of Z_k (lengths of cycles are preserved with their row sets, so the length condition is invariant along with the rest). Legality: Z_k meets $\{1, 2\}$ in 0 or 2 rows and $\{q_k, c_k\} \cap \{p, p'\} = \emptyset$, so the type of rows 1, 2 and the mark are preserved (Lemma 6.2); columns p, p' are untouched, so the frame and the ladder are preserved. For $k \neq l$ in U the column pairs $\{q_k, c_k\}, \{q_l, c_l\} \subseteq C_l$ are disjoint, so the turns act on disjoint cells and commute, and the action is free. Invariance of the data of a pair $l \neq k$ under T_k : rows $x_l, y_l \notin Z_k$, so ρ_l , its status, μ_l and all its candidates are unchanged; the goodness of candidate (l, t) depends on the column cycle of (q_l^t, c_l^t) through x_l , i.e. on the contents of columns q_l^t, c_l^t , which are not in $\{q_k, c_k\}$ for $t \leq t_0(l)$ (resp. for all candidates if $t_0(l)$ does not exist) by the defining condition of U ; hence $t_0(l)$ and C_l are unchanged. Invariance of the data of k itself: by Lemma 6.5, μ_k and the candidates $t \leq t_0(k)$ are unchanged; the candidates $t < t_0(k)$ have columns different from q_k, c_k (distinct points of the two paths), so their cycles and non-goodness are unchanged, and Z_k itself keeps its row set, so $t_0(k)$ is still good; C_k is unchanged. If $Z_k \ni 1, 2$, rows 1, 2 change at q_k, c_k , which affects none of the above. Hence U is invariant, every orbit of the cube is a cube of size $2^{|U|}$ on which the status of pair $k \in U$ is $\text{bit}_k(L) \oplus \varepsilon_k$, and $\text{Flip}_I = \emptyset$ forces one value of $(\varepsilon_k)_{k \in U}$; X is uniform. $\square$

Hypothesis 6.7 (short clean cycles at the ladder rows). *There are $T, \ell_1 = O(\log^2 n)$ and a constant $\kappa_2 > 0$ such that, uniformly in λ, α, β , with $\ell_0 = n/(16 \log^2 n)$, $M = \lfloor n/\log^4 n \rfloor$ and U as in Proposition 6.6:*

- (i) *for the sub-ladder I of the first $\min(K_0 - 1, \lfloor \ell_0 \rfloor)$ ladder pairs, $\mathbb{E}_X[2^{-|U|}; K_0 \geq \ell_0] = o(1/\log n)$;*

(ii) for a B -instance, with $\ell := d_{12} < \ell_0$ and $|C_1| \geq n/2 + \ell$: let $S = \{2\ell t : 0 \leq t \leq (n/2 - \ell)/(2\ell)\}$ (offsets at gaps 2ℓ , so that the x -rows of the diagonals D_c , $c \in S$, are pairwise distinct; $|S| \geq n/(8\ell) \geq 2\log^2 n$), let S' be the first $m_\ell := \min(|S|, \lceil n/(\log^2 n \cdot \min(\ell - 1, M)) \rceil)$ offsets of S , let $D_c^{(M)}$ be the first $\min(\ell - 1, M)$ pairs of D_c , and take U for the family $\mathcal{F}_\ell = \bigcup_{c \in S'} D_c^{(M)}$ (which has at most $2n/\log^2 n$ pairs). Then the event that fewer than $\kappa_2 \log^2 n$ of the diagonals $D_c^{(M)}$, $c \in S'$, meet U has probability $o(1/\log n)$ (one event, not one per ℓ); and symmetrically for d_{21} .

The scaling in (ii) is forced: a diagonal of $\ell - 1$ pairs meets U with probability $\approx \min(\ell - 1, M) T \ell_1 / n$, so for short arcs only a small fraction of the diagonals can be expected to meet U , and the family must be kept to $O(n/\log^2 n)$ pairs for the cleanliness and collision conditions of Proposition 6.6 to have constant cost; $|S'| \min(\ell - 1, M) \approx n/\log^2 n$ makes the expected number of diagonals meeting U of order $\log^2 n$ for every ℓ .

Proposition 6.8. *Each of Hypotheses 6.4 and 6.7 implies (L), hence the weak CGW conjecture.*

Proof. By Theorem 4.2, $|\mathbb{P}[B] - \mathbb{P}[A]| \leq \mathbb{P}[\text{Flip} = \emptyset] \leq \mathbb{P}[K_0 \geq \ell_0, \text{Flip} = \emptyset] + \mathbb{P}[A, K_0 < \ell_0] + \mathbb{P}[B, K_0 < \ell_0]$, with ℓ_0 as in the hypothesis used. *Long ladders.* K_0 is a function of the frame, hence orbit-invariant, and by Proposition 6.3 resp. Proposition 6.6 $\mathbb{P}[\text{Flip}_I = \emptyset \mid \text{orbit}] \leq e^{-Y_I/2}$ resp. $2^{-|U|}$, so $\mathbb{P}[K_0 \geq \ell_0, \text{Flip} = \emptyset] = o(1/\log n)$ by (i). *Short ladders on A.* $\mathbb{P}[A, K_0 < \ell_0] \leq 72\ell_0/n$ by Proposition 4.6. *Short arcs on B.* Fix $\ell < \ell_0$ and consider B -instances with $d_{12} = \ell$; the case $d_{21} = \ell$ reduces to it by the involution exchanging rows 1, 2, as in Proposition 5.2. The part $|C_1| < n/2 + \ell$ costs at most $8/n$ by Proposition 4.7. On $|C_1| \geq n/2 + \ell$ the family of (ii) is as required by the propositions: its x -rows lie in the disjoint depth intervals $[c + 1, c + \ell - 1]$, $c \in S$, on the frame arc from 2 to 1, and its y -rows $\pi^{-i}(2)$ on the arc from 1 to 2. Given the orbit the statuses of all pairs are independent, and the diagonals' flippabilities are independent (they are disjoint sets of pairs). Under Hypothesis 6.4, a diagonal with $Y_{D_c} \geq \kappa_0(\ell - 1)$ is flippable with conditional probability $\geq 1 - e^{-\kappa_0(\ell - 1)/2} \geq 4\theta_0$, $\theta_0 := \frac{1}{4}(1 - e^{-\kappa_0/2})$; off the exceptional event of (ii) at least $\kappa_0|S|$ diagonals qualify, so the number of flippable offsets has conditional mean $\geq 4\theta_0\kappa_0|S|$ and the fraction θ_S is $\geq 2\theta_0\kappa_0 =: \theta'_0$ except with conditional probability $\leq e^{-\theta_0\kappa_0|S|/2} \leq e^{-\theta_0\kappa_0 \log^2 n/16}$ (as $|S| \geq n/(8\ell_0) = \log^2 n/8$). Under Hypothesis 6.7, a diagonal $D_c^{(M)}$ meeting U is flippable with conditional probability $\geq \frac{1}{2}$ (the status of a pair $k \in U$ is $\text{bit}_k \oplus \varepsilon_k$ with ε uniform on the orbit); off the exceptional event of (ii) at least $\kappa_2 \log^2 n$ diagonals of S' meet U , so the number of flippable offsets in S' has conditional mean $\mu \geq \frac{1}{2}\kappa_2 \log^2 n$ and is $\geq \mu/2$ except with conditional probability $\leq e^{-\mu/8} \leq e^{-\kappa_2 \log^2 n/16}$, hence the fraction $\theta_{S'}$ is $\geq \theta_0(\ell) := \kappa_2 \log^2 n/(4m_\ell)$. In both cases the conditional failure probabilities are summable over $\ell < \ell_0$ to $o(1/n)$. The derivation of (4) from (3) applies to the average over any set of offsets determined by $\ell = d_{12}$ alone, such as S and S' , so

$$\sum_{\ell < \ell_0} \mathbb{P}[B, d_{12} = \ell, |C_1| \geq n/2 + \ell, \theta_S \geq \theta_0(\ell)] \leq \sum_{\ell < \ell_0} \frac{36}{(n-2)\theta_0(\ell)}$$

(with S' and $\theta_{S'}$ in the adaptive case), which is $36\ell_0/((n-2)\theta'_0) = O(1/\log^2 n)$ for Hypothesis 6.4,

and for Hypothesis 6.7, with $m_\ell \leq n/(\log^2 n \min(\ell - 1, M)) + 1$,

$$\begin{aligned} \sum_{\ell < \ell_0} \frac{144 m_\ell}{(n-2)\kappa_2 \log^2 n} &\leq \frac{144}{\kappa_2 \log^4 n} \left(\sum_{\ell-1 \leq M} \frac{1}{\ell-1} + \sum_{M < \ell-1 < \ell_0} \frac{1}{M} \right) + \frac{144 \ell_0}{(n-2)\kappa_2 \log^2 n} \\ &= O\left(\frac{\log n}{\log^4 n} + \frac{\ell_0}{M \log^4 n} + \frac{1}{\log^4 n}\right) = O\left(\frac{1}{\log^2 n}\right). \end{aligned}$$

The exceptional events of (ii) over all ℓ form a single event of probability $o(1/\log n)$ ($\ell = d_{12}$ is a frame statistic). Altogether $\mathbb{P}[\text{Flip} = \emptyset] = o(1/\log n)$. $\square$

What the hypothesis says, and what it costs. A ladder pair enters U when one of its T candidate column pairs has a short column cycle through x_k that is clean — a configuration of $O(\log^2 n)$ cells ($4t$ cells of rows x_k, y_k to locate the candidate, and the cycle itself). Heuristically a candidate's column cycle has length $\leq \ell_1$ with probability $\approx \ell_1/n$ (the data of §(d) show the lengths spread like a uniform variable on $\{2, \dots, n\}$), is clean with probability $\approx e^{-2s\ell_1/n}$ for a family of s pairs, and the chosen pair avoids the $\approx 2Ts$ candidate columns of the other pairs with probability $\approx (1 - 2Ts/n)^2$, so

$$\mathbb{E}|U| \approx s \cdot \frac{T\ell_1}{n} e^{-2s\ell_1/n} \left(1 - \frac{2Ts}{n}\right)^2 \asymp \log^2 n \quad \text{for } s = \frac{n}{16 \log^2 n}, T = \frac{\log^2 n}{4}, \ell_1 = \frac{\log^2 n}{2},$$

against the $C \log \log n$ ($C > 1/\log 2$) that (i) needs; the constant is small ($\approx 1/145$), so the heuristic excess over the threshold is a large- n statement. Compared with the fixed matching of (c), which offers $\approx n/2$ coordinates per pair of which a constant number are short and clean and a fraction ≈ 0.19 of those separated, the adaptive rule offers only $T = O(\log^2 n)$ coordinates per pair, all separated: the genericity constant is gone, and in exchange membership in U is a rare event per pair, so (i) is a lower-tail statement about rare local events at rows selected by the square — of the same shape as a statement that some ladder pair has a short ρ_{x_k, y_k} -cycle through p avoiding p' (which would make it flippable outright), now with a column cycle of a free pair through x_k in its place. There is a second regime: with $s \approx T \approx \ell_1 \approx \sqrt{n}$ a pair enters U with probability bounded below and $\mathbb{E}|U| \asymp \sqrt{n}$, so (i) becomes a second-moment statement about $\sqrt{n}$ events of constant probability, at the price of the arcs of ρ_k at p, p' exceeding $\sqrt{n}$ for a constant fraction of the pairs. In both regimes the lengths of the arcs of ρ_k at p, p' — which the fixed-matching version needed to be macroscopic for a positive orbit weight — enter only through the number of candidates $\min(T, \mu_k - 1)$. The collision rule in U (the chosen pair must avoid the candidate columns of *all* other pairs) is what makes U orbit-invariant — a turn can shrink the candidate set of an excluded pair, so the obvious greedy refinement is not invariant — and it is crude: it is why the asymptotic regime $2Ts/n \leq \frac{1}{2}$ is out of reach of sampling at $n \leq 100$.

Comparison with the witness lemma. Hypothesis 5.1 asks for *one* witness on a long ladder (a short ρ_{x_k, y_k} -cycle through p avoiding p' , which makes the pair flippable with no further argument), at a per-pair rate $\approx \ell'/n$. Hypothesis 6.7 asks for $C \log \log n$ good candidates at a per-pair rate $\approx T\ell_1/n$, and Hypothesis 6.4 asks for a genericity constant. Neither orbit hypothesis is implied by the witness lemma or implies it (we see no implication in either direction); the witness lemma is the least demanding heuristically and is the reference hypothesis of this paper. What the orbit method contributes is the mechanism — exact fairness compounded over commuting

involutions into exponential decay — and, with the adaptive coordinates, the observation that the mechanism needs only rare local witnesses of a second kind (column cycles of a free pair through x_k) rather than a constant; at the first ladder pair alone, (e) below, the rarity disappears too, and what is left, after §7, is a constant conditional hazard among the first $O(\log \log n)$ column pairs along the row permutation of that pair. All of these hypotheses share the same unproved core, stated in §9.

(d) Data. All scripts and logs are in the repository [14], under `s13_*.py` and `runs/s13/orbit/`. *Adaptive coordinates.* On instances with a uniformly random mark at $n = 30, 50, 100$ (family the first $\min(K_0 - 1, n/8)$ ladder pairs, $T = 6$, $\ell_1 = 6, 8, 10$; 2400, 3000, 900 instances), for a random $k \in U$ the turn of Z_k was applied and checked to toggle the status of (x_k, y_k) , to leave the status, μ_l , $t_0(l)$, C_l of every other pair and the set U unchanged, and to preserve the type of rows 1, 2 and the frame (682, 697, 215 turns, no failure; 444 and 160 more with families of 12 and 30 pairs); the bound $\mathbb{P}[\text{Flip}_I = \emptyset] \leq \mathbb{E}[2^{-|U|}]$ holds in every K_0 class. At these sizes 30%, 23%, 16% of the ladder pairs have a good candidate and 13%, 5%, 3% are in U (8% and 1.4% with the family of 30 pairs at $n = 100$): $2Ts \geq n$ there, so the collision rule removes most pairs. *Column-cycle lengths and the fixed-matching coordinates.* The length of the column cycle of a free pair through a ladder row is spread like a uniform variable on $\{2, \dots, n\}$ (about $1/n$ per length; the turn is legal for 98% of the 2-cycles and 64% of the long cycles, and avoids y for 97% resp. 27% of them, at $n = 100$). For the fixed matching of (c): $\mathbb{P}[(q, c) \text{ separated} \mid \text{clean, length } \ell] = 0.193\text{--}0.201$ for $\ell = 2, \dots, 8$ against 0.195 over all column pairs ($n = 100$), i.e. no detectable dependence on the cycle length; with $|\mathcal{T}| \approx 42$ coordinates the orbit average of the parallel bit is 0.478 from a crossed and 0.523 from a parallel start ($n = 100$); with the fixed matching of consecutive free columns $|\mathcal{T}| = 1.7$ resp. 3.0 clean coordinates per pair ($n = 50, 100$) and $\mathbb{E}[\bar{y}] = 0.159$ resp. 0.189; and the orbit weights of different ladder pairs are uncorrelated to within sampling error ($\text{Var}(Y_I)/(|I|\text{Var } \bar{y}) = 0.99$ and 1.14 at $n = 50, 100$, the excess coming from the variation of $|I|$). The ladder data of Remark 4.4 complete the picture: $\mathbb{P}[\text{Flip} = \emptyset] \approx 4/n$, as if the ladder coins were fair and independent.

(e) The first ladder pair alone. The adaptive coordinates of (c') specialise to the single pair $(x_1, y_1) = (\pi^{-1}(1), \pi^{-1}(2))$ with none of the side conditions of Proposition 6.6 (there are no other pairs to protect), and the resulting involution combines with the first ladder trade T_1 into a bound on $\mathbb{P}_X[B] - \mathbb{P}_X[A]$ by an event selected by that one pair. Since $\alpha, \beta \geq 2$ the pair always consists of rows outside $\{1, 2\}$: $x_1 = 2$ would mean $p' = \sigma(p)$ and $y_1 = 1$ would mean $p = \sigma(p')$ (so $K_0 \geq 2$ always). Write $\rho = \rho_{x_1, y_1}$, F for the event that (x_1, y_1) is flippable (parallel), and ν for the minimum of the two arc lengths from p to p' and from p' to p (crossed case) or of $|Q_p|, |Q_{p'}|$ (parallel case) — the quantity written μ_{x_1, y_1} in (c'); we use ν here because μ is the split type elsewhere in the paper.

Lemma 6.9. $\nu \geq 2$ always.

Proof. $\rho(p) = p'$ would mean $L(y_1, p') = L(x_1, p) = L(1, p')$, impossible in a column; $\rho(p') = p$ would mean $L(x_1, p') = L(y_1, p) = L(2, p')$, again impossible in a column; and ρ has no fixed points since $x_1 \neq y_1$. $\square$

For $1 \leq t < \nu$ the column pair $\{q_t, c_t\} = \{\rho^t(p), \rho^t(p')\}$ (written q^t, c^t in (c')) is free, consists of two distinct columns, and is separated by the mark, as observed in (c'). Let $\tau_t = \delta_{q_t, c_t}$ be the permutation of rows induced by the column pair ($\tau_t(r)$ is the row of $L(r, q_t)$ in column c_t) and

$Z_t(r)$ its cycle through row r . Call t *good* if x_1 and y_1 lie in different cycles of τ_t and at least one of $Z_t(x_1)$, $Z_t(y_1)$ meets $\{1, 2\}$ in 0 or 2 rows (so that its turn is legal); let

$$G = \{\text{some } t < \nu \text{ is good}\},$$

and on G let t_0 be the least good t and Z the admissible cycle at t_0 (the one through x_1 if it is admissible, otherwise the one through y_1).

Proposition 6.10 (single-pair involution). *Turning Z (exchanging the entries of columns q_{t_0}, c_{t_0} in the rows of Z) is an involution of $X \cap G$ that preserves the frame π (hence x_1, y_1 , A and B), toggles F , and fixes t_0 and Z . Consequently $\mathbb{P}_X[F \cap G] = \mathbb{P}_X[F^c \cap G]$ and $|\mathbb{P}_X[F] - \frac{1}{2}| \leq \frac{1}{2}\mathbb{P}_X[G^c]$.*

Proof. Columns p, p' are untouched, so π , x_1, y_1 , the mark and the rows outside Z are unchanged. Rows 1, 2 are both in Z or both outside; if both are in, they become $L_1 \circ (qc)$ and $L_2 \circ (qc)$ (row 1 is then no longer the identity, which the unnormalised definition of X allows), so the new row permutation is $(qc)\sigma(qc)$, of the type of σ , and as $q, c \notin \{p, p'\}$ it still has p' at distance α from p : the turned square is in X . Exactly one of x_1, y_1 is in Z (they lie in different τ_{t_0} -cycles), so $\rho' = \rho \circ (qc)$ if $x_1 \in Z$ and $\rho' = (qc) \circ \rho$ if $y_1 \in Z$. In the first case Lemma 6.5 applies with $t = t_0$. In the second, $(qc) \circ \rho = \rho \circ (u' u)$ with $u = \rho^{t_0-1}(p)$, $u' = \rho^{t_0-1}(p')$ (possibly $t_0 = 1$ and $\{u, u'\} = \{p, p'\}$), and the computation in the proof of Lemma 6.5, read with $t_0 - 1$ in place of t (for $t_0 = 1$ the paths are empty and the computation is immediate), shows that the ρ -paths from p to u and from p' to u' are unchanged, that $\rho'(u) = c$ and $\rho'(u') = q$ (the ordered pair at step t_0 is reversed, the unordered pair is $\{q, c\}$), and that crossed with arcs (a, b) becomes parallel with cycle lengths (b, a) and conversely. In both cases $\nu' = \nu$, the candidate pairs are the same for $t \leq t_0$, and F toggles. The column pairs with $t < t_0$ are disjoint from $\{q, c\}$, so their τ_t , hence their badness, are unchanged. At t_0 the turn replaces τ_{t_0} by τ_{t_0} with the cycle Z reversed (or by the inverse of that, if the ordered pair is reversed), so x_1, y_1 are still in different cycles, Z is still a cycle with the same intersection with $\{1, 2\}$, and the other cycle through $\{x_1, y_1\}$ is unchanged: t_0 and Z are the same in the turned square, and the turn is an involution. The identity follows, and $\mathbb{P}[F] \geq \frac{1}{2}\mathbb{P}[G]$, $\mathbb{P}[F^c] \geq \frac{1}{2}\mathbb{P}[G]$ give the bound. $\square$

Proposition 6.11 (the first ladder pair controls the bias). *For every $n, \lambda, \alpha, \beta$,*

$$\mathbb{P}_X[B] - \mathbb{P}_X[A] = \mathbb{E}_X[(\mathbf{1}_B - \mathbf{1}_A)(\mathbf{1}_{F^c} - \mathbf{1}_F); G^c];$$

consequently $|\mathbb{P}_X[B] - \frac{1}{2}| \leq \frac{1}{2}\mathbb{P}_X[G^c]$, $\mathbb{P}_X[A], \mathbb{P}_X[B] \geq \frac{1}{2}\mathbb{P}_X[G]$, and $|\bar{q}(\mu \rightarrow \lambda) - 1| \leq 2\mathbb{P}_X[G^c]/(1 - \mathbb{P}_X[G^c])$.

Proof. Write $B_{FG} = \mathbb{P}_X[B \cap F \cap G]$ and so on. The trade T_1 of Theorem 4.2 ($k = 1$) is an involution of $X \cap F$ whose frame is $(x_1 y_1) \circ \pi$, so it exchanges A and B : $B_{FG} + B_{FG^c} = A_{FG} + A_{FG^c}$. The turn of Proposition 6.10 is an involution of $X \cap G$ exchanging F and F^c and preserving π , hence A and B , which are functions of π : $B_{FG} = B_{F^cG}$ and $A_{FG} = A_{F^cG}$. Hence

$$\begin{aligned} \mathbb{P}[B] - \mathbb{P}[A] &= (B_{F^cG} + B_{FG^c}) - (A_{F^cG} + A_{FG^c}) = (B_{FG} - A_{FG}) + (B_{F^cG^c} - A_{F^cG^c}) \\ &= (A_{FG^c} - B_{FG^c}) + (B_{F^cG^c} - A_{F^cG^c}), \end{aligned}$$

which is the displayed expectation; its absolute value is at most $\mathbb{P}[G^c]$. (T_1 need not preserve G .) The last inequality is $|\bar{q} - 1| = |\mathbb{P}[B] - \mathbb{P}[A]|/\mathbb{P}[A]$ with $\mathbb{P}[A] \geq \frac{1}{2}(1 - \mathbb{P}[G^c])$. $\square$

Hypothesis 6.12 (first-pair hypothesis). $\sup_{\lambda, \alpha, \beta} \mathbb{P}_{X(n, \lambda; \alpha, \beta)}[G^c] = o(1/\log n)$.

Corollary 6.13. *Hypothesis 6.12 implies Conjecture 1.1, with $d_{TV}(P_n, Q_n) = O(\log n \cdot \sup \mathbb{P}_X[G^c] + n^{-1+o(1)})$.*

Proof. With $s = \sup \mathbb{P}_X[G^c]$, Proposition 6.11 gives $|\bar{q}(\mu \rightarrow \lambda) - 1| \leq 2s/(1-s)$ for every adjacent pair, i.e. EQ(δ) with $\delta = 2s/(1-s) = o(1/\log n)$, and Theorem 3.3 applies. $\square$

What the first-pair hypothesis asks, and what it costs. The event G^c is selected by the first ladder pair but is not local to it: it involves the column cycles of the pairs $\{q_t, c_t\}$ through all rows, and rows 1, 2 through admissibility. It is again a lower tail of a count of events at square-selected positions — the obstacle is unchanged in kind — and it decomposes as

$$\mathbb{P}_X[G^c] \leq \mathbb{P}_X[\nu \leq M] + \mathbb{P}_X[\nu > M, t = 1, \dots, M \text{ all bad}].$$

The first term is an upper bound on a short arc or cycle of ρ_{x_1, y_1} at p, p' ; before §7 nothing of the kind was available at a frame-selected pair (Proposition 4.7 bounds a different frame event); heuristically it is $\approx 4M/n$, and §7 proves $O(\Delta^{2M}/n)$. The second term — among the first M candidates, which lie on disjoint free column pairs, not all are bad — is a conditional statement about M events of constant probability, and with $M = C \log \log n$ it needs only a hazard bounded below: if $\mathbb{P}_X[t \text{ bad} \mid \nu > t, 1, \dots, t-1 \text{ bad}] \leq 1-c$ for $t \leq M$ then the second term is $\leq (1-c)^M = (\log n)^{-C \log(1/(1-c))}$, which is $o(1/\log n)$ for C large (in the data a candidate is bad with conditional probability ≈ 0.6 , i.e. good with ≈ 0.4 , Remark 6.14). Under such a hazard bound the failure of G is dominated by short arcs or cycles at the mark, not by a large count of events: the first-pair hypothesis is then the ν -tail at the selected pair together with a constant conditional hazard over $O(\log \log n)$ candidates. Compared with Hypothesis 5.1, at the level of heuristics and data only: one row pair instead of $\geq n/\log^2 n$; per-event probability ≈ 0.4 (the random-permutation value for “ $x_1 \not\sim y_1$ under τ_t and one of the two cycles has even intersection with $\{1, 2\}$ ”) instead of ℓ'/n ; $O(\log \log n)$ events instead of $\geq \log^2 n$ ($\log^3 n$ at the parameters of Proposition 5.2); no offset families, no splice bound, no conditioning on K_0 ; and the quantity hypothesised is one the data measure directly ($n\mathbb{P}_X[G^c]$ is flat at 5.5 for $30 \leq n \leq 150$), whereas the threshold $\ell' = \log^5 n$ of Hypothesis 5.1 exceeds every n we can sample. Against this: the candidates are generated by one permutation ρ_{x_1, y_1} and are dependent (the measured dependence is positive, the unfavourable direction); and admissibility couples G to rows 1, 2. The ν -tail, which is the dominant term in the data, is proved to be $O(\Delta^{2M}/n)$ in §7 below (Corollary 7.3), so that only the hazard term remains (Hypothesis 7.4). Neither of Hypotheses 6.12 and 5.1 is known to imply the other: $\mathbb{P}_X[F^c] \geq \frac{1}{2}\mathbb{P}_X[G]$, so the first pair alone cannot give (L), and Hypothesis 5.1 says nothing about G . So Hypothesis 6.12 is a third alternative sufficient condition, the simplest to state and the only one whose hypothesised quantity we can measure; the reference hypothesis of the paper remains Hypothesis 5.1. The minimal instance (7) of §9 follows from Hypothesis 6.12 (with $c = \frac{1}{2} - o(1)$) but not evidently from Hypothesis 5.1, and is sufficient for neither.

The staircase events. The event $\{\nu \leq M\}$ is the union over $2 \leq k \leq M$ of the four events $E_k(a, b)$, $a, b \in \{p, p'\}$: the ρ -path $w_0 = a, w_1, \dots$ (with $w_{i+1} = \rho(w_i)$, i.e. $L(y_1, w_{i+1}) = L(x_1, w_i)$) first meets $\{p, p'\}$ at step k , at $w_k = b$. Each is a “staircase” of $2k$ cells in rows x_1, y_1 , (x_1, w_i, t_i) and (y_1, w_{i+1}, t_i) for $0 \leq i < k$ with $t_i = L(x_1, w_i)$ and $w_1, \dots, w_{k-1} \notin \{p, p'\}$ distinct, together with the two anchor cells $L(x_1, p) = L(1, p')$, $L(y_1, p) = L(2, p')$ that select the pair; all of them lie in the two rows $x_1, y_1 \notin \{1, 2\}$. By the exchangeability of rows $3, \dots, n$,

$\mathbb{P}_X[E_k(a, b)] = (n-2)(n-3)\mathbb{P}_X[E_k(a, b), (x_1, y_1) = (3, 4)]$, and the latter event is a union, over at most n^{2k-1} choices of the free columns and symbols, of events $\{L \supseteq Q\}$ with Q a partial Latin square of at most $2k+2$ cells in rows 3, 4 (for $(a, b) = (p, p)$, mixed, (p', p') : $2k, 2k+1, 2k+2$ cells and $n^{2k-3}, n^{2k-2}, n^{2k-1}$ parameters). So an upper bound on the ν -tail follows from a *fixed-position* bound in completions of a fixed rectangle,

$$\mathbb{P}[Q \subseteq L \mid L \supseteq R] \leq (C_0/n)^{|Q|} \quad \text{for partial Latin squares } Q \text{ with } |Q| \leq 2M+2 \text{ cells in two rows,} \quad (6)$$

L a uniform completion of the fixed $2 \times n$ rectangle R (each $S(R)$ contributes to $\mathbb{P}_X$ with the same weight, rectangles of one type having equal completion counts): (6) gives $\mathbb{P}_X[E_k(a, b)] \leq (n-2)(n-3)n^{2k-3+t}(C_0/n)^{2k+t} \leq C_0^{2k+2}/n$, $t \in \{0, 1, 2\}$, hence, summing the four events for each k and the geometric series over $2 \leq k \leq M$, $\mathbb{P}_X[\nu \leq M] \leq 8C_0^{2M+2}/n$ for $C_0 \geq 2$. The always-legal row trades give (6) only to one factor $1/n$ per row (an elementary rigid switching along the staircase gives $\mathbb{P}_X[\nu \leq M, f_{\text{adm}} \geq \theta n] \leq 4M/(\theta n)$ with f_{adm} the number of admissible rows, which leaves a lower tail; [14]); two cells in one row need the η -switching of Allsop and Morris [1], whose per-cell bound we extend to partial Latin squares containing two complete rows in §7. That section proves (6) with an absolute constant $C_0 = \Delta \leq 66$.

Remark 6.14 (data: completions of a fixed rectangle). To test the hypotheses on types that uniform sampling never produces we ran the Jacobson–Matthews chain restricted to the completions of a fixed $2 \times n$ rectangle (`jmfix.c`: moves touching rows 1, 2 are never proposed from a proper state and are rejected from an improper one; the walk is a simple random walk on the restricted move graph, every proper state having restricted degree $(n-2)n(n-3)$, so its stationary law is uniform on the proper states of the component of the start). Connectivity of the restricted graph is not proved; at $n = 5, 6$ (types (5), (6), (2, 2, 2); 36, 4032, 5376 completions, enumerated exactly) every completion is reached from two different starts with χ^2 statistics within one standard deviation of uniformity, and at $n = 7$ the sampled $\mathbb{P}[B \mid R] = 0.4965$ agrees with the exact 0.497076 over all 6566400 completions of $R = (\text{id}, (1\ 2 \cdots 7))$. Mixing is not analysed; runs with different thinning and starts agree within noise. Any two rectangles of type λ are isotopic by a column permutation and a symbol relabelling fixing rows 1, 2, which preserve completion counts and carry a mark of class (m, α) to one of the same class; every event considered here is defined through the cycle structures of row and column permutations and is invariant under such isotopies. Hence $\mathbb{P}_X$ of such an event equals its probability under a uniform completion of any one rectangle R of type λ with any one mark of the class, and averaging over the marks of the class in R (which are equivalent under the isotopies preserving R) only reduces the variance. Findings (assuming connectivity and mixing of the sampler) ($n = 30$: $6 \cdot 10^4$ squares per type; $n = 50$: $4 \cdot 10^4$; $n = 100$: $1.2 \cdot 10^4$; types (n) , $(n/2, n/2)$, $(n-2, 2)$, $(10, 10, 10)$, $(4, 2^{13})$, $(6, 3^8)$, $(4, 2^{23})$, $(6, 3^{14}, 2)$, $(4, 2^{48})$; marks at $\alpha = 2$, $\approx m/3$, $\lfloor m/2 \rfloor$): (i) the point estimates of $\mathbb{P}_X[B]$ lie in $[0.4974, 0.5024]$ in every class at $n \leq 50$ (each with sampling error ≈ 0.0025 , so the spread is noise), i.e. by Theorem 2.4 the point estimates of $C_n(\lambda)/C_n(\mu)$ lie in $[0.995, 1.005]$ (each ± 0.005) for every tested adjacent pair, including $(4, 2^{13}) \rightarrow (2^{15})$ and $(6, 3^8) \rightarrow (3^{10})$, types of P_n -probability $e^{-\Theta(n)}$; $[0.495, 0.504]$ at $n = 100$ (error ≈ 0.005). (ii) $\mathbb{P}_X[F] \in [0.4988, 0.5017]$ in every class at $n \leq 50$ (0.4896 at $n = 7$, so $\frac{1}{2}$ is not an identity). (iii) Along the ladder, $\mathbb{P}[\text{no witness among the first } K \text{ pairs} \mid K_0 - 1 \geq K]$ equals $(1 - w_1)^K$ within sampling error for $K \leq 12, 20, 40$ at $n = 30, 50, 100$ and thresholds $\ell' = 4, 8, 16, 24$, where w_1 is the first-pair witness rate (0.100, 0.060, 0.030 at $\ell' = 4$); the hazard $\mathbb{P}[\text{witness at } k \mid \text{none before } k]$ is flat in k to three digits; $\mathbb{P}[\text{none of the first } K \text{ ladder pairs flippable} \mid K_0 - 1 \geq K] = 2^{-K}$ within noise to $K = 12\text{--}15$ ($7 \cdot 10^5$ instances at $n = 30$); across the diagonals D_c of B -instances

with $d_{12} = \ell \leq n/4$, which share their y -rows, the number of diagonals carrying a witness vanishes with the binomial probability and has binomial lower tails, for all $\ell \leq 12$. All of this is identical for the extreme types. (iv) For Hypothesis 6.12, with all candidates $t < \nu$: $n\mathbb{P}_X[G^c] = 5.5, 5.5, 5.5, 5.4$ at $n = 30, 50, 100, 150$ (type (n) , $\alpha = 2$; 5.6, 5.5 for $\alpha = n/2$ at $n = 100, 150$; 5.5 for $(4, 2^{73})$ at $n = 150$; 0.61 at $n = 7$); $\mathbb{E}\nu = 6.3, 9.6, 18$ at $n = 30, 50, 100$ and $\mathbb{P}[\nu = k] \approx 4/n$ for small k at every n ; $\mathbb{P}[G^c \mid \nu = k] = 0.60, 0.30, 0.20, 0.14, 0.07$ for $k = 2, \dots, 6$ at $n = 150$ against $0.59^{k-1} = 0.59, 0.35, 0.21, 0.12, 0.07$, the same at every n and type (a slight positive dependence among the candidates, which share the rows x_1, y_1); $t = 1$ is good with probability 0.41 in every class at $n \leq 100$. Other types were measured only with $t \leq 8$: $\mathbb{P}[\text{some } t \leq 8 \text{ good}] = 0.81, 0.88, 0.94$ at $n = 30, 50, 100$ in every class. The identity of Proposition 6.11 and its two ingredients hold to all digits on the exact enumerations at $n = 5, 6$ (every class) and within sampling error at $n = 7$ (-0.0052 and -0.0058 , each ± 0.001 , against the exact -0.005848). None of these data reach the regime of any hypothesis ($\ell' \geq \log^5 n$, failure $o(1/\log n)$); they exclude a visible dependence between the events, or a visible effect of the type, at the first- and second-moment level for $n \leq 100$.

7 The fixed-cell bound in completions of a rectangle

Throughout this section R is a fixed $2 \times n$ Latin rectangle occupying rows 1, 2, Q is a partial Latin square with cells in rows other than 1, 2 such that $R \cup Q$ is a partial Latin square, and L is uniform on $\{L \in \mathcal{L}_n : L \supseteq R \cup Q\}$; we write $S(R)$ for the completions of R . Nothing in this section moves rows 1, 2, and the mark plays no role. (In this section α, β denote the density parameters of [1], not the mark parameters; the mark parameters return in Corollary 7.3. Their statistic $\nu(L)$ is not used; our ν is the arc minimum of §6(e).) We use the switchings of Allsop and Morris [1] and follow their notation (§§2–4 there): $\rho_L(r, r', c)$ is the row cycle of rows r, r' through column c , $\gamma_L(c, c', r)$ the column cycle of columns c, c' through row r , $\zeta_L(c, c', r, r')$ the cross-switch, and $\eta_L(r, r', c)$ the η -switch, which replaces the entries $(r, c, s), (r', c, s'), (r, c', s')$ — where $s = L[r, c]$, $s' = L[r', c]$ and c' is the column with $L[r, c'] = s'$ — by $(r, c, s'), (r', c, s), (r, c', s)$ and otherwise changes only rows r' and r'' , the row with $L[r'', c'] = s$; it is of *Type One* if $r'' = r'$ (then $\rho_L(r, r', c)$ is an intercalate), and by their Lemma 2.8 a square has at most n preimages under an η -switch of Type One and at most one otherwise. Their Theorem 1.1 bounds $\mathbb{P}[L \supseteq P]$ for partial Latin squares P occupying at most αn rows and βn columns with $2\alpha + \beta < 1$, and so does not apply to $P = R \cup Q$, which occupies every column. Their proof does, with changes we describe in full. Call rows 1, 2 *special* and write $E = \{1, 2\} \cup R_Q$.

Theorem 7.1 ([1], Theorem 3.1, with two complete rows). *Let R, Q be as above with $|R_Q| + 6 \leq \alpha n$, $|C_Q| \leq \beta n$ and $2\alpha + \beta < 1$. Let $r, r' \notin \{1, 2\}$ be distinct rows and c a column, and suppose that Q has no cell in row r' outside column c . Then, for $n \geq n_0(\alpha, \beta)$,*

$$\mathbb{P}[\eta_L(r, r', c) \text{ is of Type One} \mid L \supseteq R \cup Q] \leq \frac{D}{n}, \quad D = D(\alpha, \beta) \rightarrow 62 \text{ as } \alpha, \beta \rightarrow 0.$$

The constant is not optimised. The proof is that of [1] with three changes: the row-exclusion set R_P becomes E and the column-exclusion set C_P becomes C_Q (the cells of R need no exclusion because no switch below changes rows 1, 2); the cross-switches of their Claim 2 and the row switches of their Claim 3 are restricted to pairs of rows in a common *arc* of a column cycle, so that no column part of a switch passes through a special row; and the statistic $\nu(L)$ of their Claims 3–4 is replaced by one that also counts the rows on the (c_2, c_3) -cycles through the special

rows, so that the column-cycle switch of their Claim 4 is only ever applied to a cycle avoiding rows 1, 2.

The recursion and Claim 1. For $L \supseteq R \cup Q$ put $r_0 = r$, $c_0 = c$, $s_0 = L[r_0, c_0]$, $r_1 = r'$, and let c_1 be the column with $L[r_0, c_1] = L[r_1, c_0]$. Recursively, for $j \geq 1$, $s_j = L[r_j, c_j]$; r_{j+1} is the row with $L[r_{j+1}, c_j] = s_{j-1}$ and c_{j+1} the column with $L[r_j, c_{j+1}] = s_{j-1}$; X_j is the set of squares with $s_j = s_{j-1}$, and Y_j the set with $s_j \neq s_{j-1}$, $r_{j+1} \notin E \cup \{r_\ell : \ell \leq j\}$ and $c_{j+1} \notin C_Q \cup \{c_\ell : \ell \leq j\}$ (with Y_0 all squares); $p_{j-1} = \mathbb{P}[X_j \mid Y_{j-1}]$, so that p_0 is the probability to be bounded. Their Claim 1,

$$p_{j-1} \leq \frac{np_j + 1}{n(1 + p_j - 2\alpha - \beta) - 3j - 2} \quad (1 \leq j < n(1 - 2\alpha - \beta)/3),$$

holds verbatim. Its switch $\eta_L(r_j, x, c_j)$, for $L \in X_j$ and a row $x \notin E \cup \{r_\ell\}$ with y (the column with $L[r_j, y] = L[x, c_j]$) outside $C_Q \cup \{c_\ell\}$ and x' (the row with $L[x', y] = s_{j-1}$) outside $E \cup \{r_\ell\}$, changes only rows r_j, x, x' , all outside $\{1, 2\}$, so rows 1, 2 are untouched; the cells of Q it could touch are (r_j, y) , excluded, (r_j, c_j) , which is not a cell of Q ($c_j \notin C_Q$ for $j \geq 2$, and for $j = 1$ the only cell of Q allowed in row r_1 is in column $c_0 \neq c_1$), and cells in rows x, x' , excluded. The maps $x \mapsto y \mapsto x'$ are injective, so at most $|C_Q| + j + 1$ choices of x are lost to y and at most $|E| + j + 1$ to x' , giving at least $n - 2|E| - |C_Q| - 3j - 3 \geq n(1 - 2\alpha - \beta) - 3j - 3$ switches from each $L \in X_j$ into Y_j , and the reversibility count is theirs.

Setting for Claims 2–4. For $L \in X_3$ let $Q(L)$ consist of the cells of Q and the ten structure cells $\{(r_\ell, c_\ell), (r_{\ell+1}, c_\ell), (r_\ell, c_{\ell+1}) : \ell \leq 2\} \cup \{(r_3, c_3)\}$; a row is *free* if it is not in $E \cup \{r_0, r_1, r_2, r_3\}$, and $f \geq n - |E| - 4 \geq n(1 - \alpha)$ is the number of free rows. (The row r_0 need not be free and may lie anywhere; its cells in columns c_2, c_3 are not in $Q(L)$ and are never read by the recursion, which reads only the ten structure cells.) Consider the partition of the rows into (c_2, c_3) -column cycles. As in [1]: $\gamma_L = \gamma_L(c_2, c_3, r_1)$ is one of them; $\{r_2, r_3\}$ is another (an intercalate, since $s_3 = s_2$); and $c_2, c_3 \notin C_Q$. Call a cycle *active* if it is γ_L or contains a special row (so there are one to three active cycles, and $\{r_2, r_3\}$ is never active), let $\mathcal{S}(L)$ be the set of free rows on active cycles, and $g(L) = f - |\mathcal{S}(L)|$ the number of free rows on inactive cycles. An inactive cycle other than $\{r_2, r_3\}$ avoids rows 1, 2 and r_1, r_2, r_3 , so switching it changes no cell of $Q(L)$ or R .

Give each active cycle D (a cycle, not to be confused with the constant D) a *reference row*: r_1 if $D = \gamma_L$, otherwise row 1 if $1 \in D$, otherwise row 2. List the rows of D as $x'_1, \dots, x'_m$ in the cyclic order from the reference row x'_1 read from column c_2 (the row after x is the row whose c_2 -entry is the c_3 -entry of x , as on p. 16 of [1]), and let the free rows of D in this order be $x_1, \dots, x_d$. The special rows of D other than the reference row cut the list into *arcs*; a pair (i, j) , $i < j$, of free positions is a *same-arc pair* if no special row and no reference row lies strictly between x_i and x_j . The free rows of the active cycles are the disjoint union of the free rows of the arcs, and there are at most three arcs in all (if γ_L contains both special rows it has three arcs and is the only active cycle; if one, it has two and the other special row's cycle has one; if none, γ_L has one and the special rows lie on one cycle with two arcs or on two cycles with one each). As in [1], (L, i, j) is *good* if $\rho_L(x_i, x_j, c_3)$ does not hit column c_2 , and *bad* otherwise.

Claim 2' (good same-arc pairs, in aggregate). For $a \geq 4$ let $\mathcal{A}_a$ be the set of triples (L, D, A) with $L \in X_3$, D an active cycle of L and A an arc of D with exactly a free rows, and

$N_a = |\mathcal{A}_a|$. Then

$$\sum_{(L,D,A) \in \mathcal{A}_a} \#\{\text{good pairs } (i,j) \text{ with } x_i, x_j \in A\} \geq N_a \frac{a(a-1)(a-3)}{9(a-2)} \geq N_a \frac{(a-3)^2}{9}.$$

This is a statement about the class $\mathcal{A}_a$, as their Claim 2 is about the class $\mathcal{A}_k$; a single arc may have no good pair.

Proof. This is their proof run inside the arcs of size a . For (i,j) in A let $W_A(i,j)$ be the set of pairs (u,v) in A with exactly one of u,v strictly between i and j and $\{u,v\} \cap \{i,j\} = \emptyset$; then $|W_A(i,j)| = (j-i-1)(a-(j-i)-1)$ in arc coordinates, $\sum_{(i,j) \subseteq A} |W_A(i,j)| = a(a-1)(a-2)(a-3)/12$ and $|W_A| \leq (a-2)^2/4$ (their (5)–(6) with $n-k$ replaced by a). For a bad (L, i_0, j_0) in A and $(u,v) \in W_A(i_0, j_0)$: if (L, u, v) is good, map to (L, u, v) (Kind 1); if it is bad, cross-switch on $\zeta_L(c_2, c_3, x_u, x_v)$ (Kind 2). The cross-switch is well defined ((L, u, v) is bad, and x_u, x_v are not adjacent in the list since a position of $\{i_0, j_0\}$ separates them). It changes the rows x_u, x_v , which are free, and the entries in columns c_2, c_3 of the rows strictly between x_u and x_v in the list, which are rows of D that are neither special nor the reference row (the pair is same-arc) nor r_2, r_3 (not on D); such rows may include r_0 or rows of R_Q , whose cells in columns $c_2, c_3 \notin C_Q$ are not in $Q(L) \cup R$. So it changes no cell of R or $Q(L)$, and $L' \in X_3$ with the same r_ℓ, c_ℓ, s_ℓ . By their Lemma 2.6 the row set of D is unchanged, and by their Remark 2.7 the list of L' is that of L with the segment strictly between x_u and x_v reversed; so the special and reference rows keep their positions, A is an arc of L' with the same free rows, $(u,v) \in W_A(i', j')$ for the new positions (i', j') of x_{i_0}, x_{j_0} , and (L', i', j') is good by their Lemma 2.6(2); hence $(L', D, A) \in \mathcal{A}_a$. The inverse counts are theirs: a good (L', D', A', i', j') has at most $|W_{A'}(i', j')|$ Kind-1 preimages (the interleaving relation is symmetric) and at most $|W_{A'}(i', j')|$ Kind-2 preimages (one for each $(u,v) \in W_{A'}(i', j')$, by their Lemma 2.5 and Remark 2.7, the inverse cross-switch again lying inside A'). Summing over $\mathcal{A}_a$, $\sum_{\text{bad}} |W| \leq 2 \sum_{\text{good}} |W|$, whence $N_a a(a-1)(a-2)(a-3)/12 = \sum_{\text{all}} |W| \leq 3 \sum_{\text{good}} |W| \leq 3 \cdot \#\{\text{good}\} \cdot (a-2)^2/4$. The last inequality of the claim is $a(a-1) \geq (a-2)(a-3)$. $\square$

Claim 3' (free rows off the active cycles). $\sum_{L \in X_3} g(L) \geq f|X_3|/30$ for $f \geq 501$.

Proof. Double count the row switches $\rho_L(x_i, x_j, c_3)$ over $L \in X_3$, active cycles D of L and good same-arc pairs (i,j) on D . *Forward.* Such a switch exchanges the entries of the free rows x_i, x_j on the columns of their row cycle through c_3 , which does not contain c_2 ; it changes no cell of R or $Q(L)$, so $L' \in X_3$ with the same r_ℓ, c_ℓ, s_ℓ . The (c_2, c_3) -column permutation (the c_2 -entry of a row $\mapsto$ its c_3 -entry) changes only at rows x_i, x_j , whose c_3 -entries are exchanged: it is composed with the transposition of those two symbols, which lie on one cycle. So every cycle other than D is unchanged and D splits into two: one consists of the rows outside the segment strictly between x_i and x_j together with x_i , and contains the reference row and every special row of D ; the other, S , consists of the rows strictly between together with x_j , and contains no special row because (i,j) is a same-arc pair (this is their Lemma 2.4(1), with the orientation of our list). Hence the active cycles of L' are those of L with D replaced by $D \setminus S$ (if $D = \gamma_L$ then $\gamma_{L'} = D \setminus S \ni r_1$; otherwise $\gamma_{L'} = \gamma_L$), S is inactive in L' , $x_i \in \mathcal{S}(L')$, $x_j \in S$, and $g(L') = g(L) + |S \cap \text{free}|$. Let $F(L)$ be the number of such switches from L (distinct pairs of rows give distinct switches, and a pair of rows lies on one cycle). Summing Claim 2' over a , with $h(a) = ((a-3)_+)^2/9$,

$$\sum_{L \in X_3} F(L) \geq \sum_{a \geq 4} \sum_{(L,D,A) \in \mathcal{A}_a} \#\{\text{good pairs in } A\} \geq \sum_{a \geq 4} N_a h(a) = \sum_{L \in X_3} \sum_A h(|A|),$$

and for each L , since the arcs partition $\mathcal{S}(L)$ into at most three parts, $\sum_A h(|A|) \geq \frac{1}{27}(\sum_A (|A| - 3)_+)^2 \geq \frac{1}{27}((f - g(L) - 9)_+)^2$. *Backward.* If L' arises from (L, D, i, j) then L is obtained from L' by switching $\rho_{L'}(x_i, x_j, c_3)$ (a row cycle switch is an involution), so L is determined by the unordered pair $\{x_i, x_j\}$, of which one row lies in $\mathcal{S}(L')$ and the other on an inactive cycle of L' ; the number of preimages of L' is at most $|\mathcal{S}(L')|g(L') \leq fg(L')$. *Conclusion.* With $\bar{g} = |X_3|^{-1} \sum_L g(L)$, Jensen's inequality for the convex function $t \mapsto ((f - t - 9)_+)^2$ gives $|X_3|((f - \bar{g} - 9)_+)^2/27 \leq \sum_L \sum_A h(|A|) \leq \sum_L F(L) \leq f \sum_{L'} g(L') = f|X_3|\bar{g}$. If $\bar{g} < f/30$ then $(f - \bar{g} - 9)^2 > (29f/30 - 9)^2 \geq 0.9f^2 > 27f\bar{g}$ for $f \geq 501$, a contradiction. $\square$

Claim 4'. $p_2 \leq 60/f \leq 60/(n(1 - \alpha))$ for $f \geq 501$.

Proof. Their proof with the set of admissible rows changed. For $L \in X_3$ let x be a free row on an inactive cycle ($g(L)$ choices) and $\rho = \rho_L(r_3, x, c_3)$. If ρ does not hit column c_2 , switch on ρ to obtain L' ; otherwise first switch on the (c_2, c_3) -cycle through x to obtain L'' , then on $\rho_{L''}(r_3, x, c_3)$ to obtain L' . The cycle through x is inactive, so it avoids rows 1, 2, r_1 (on γ_L) and r_2, r_3 (the intercalate), and $c_2, c_3 \notin C_Q$: switching it changes no cell of R or $Q(L)$, so $L'' \in X_3$ with the same data; and $\rho_{L''}(r_3, x, c_3)$ does not hit c_2 , because the switch composes the row permutation of (x, r_3) with the transposition of $L[x, c_2], L[x, c_3]$, two symbols of one cycle, which splits that cycle with the two symbols, hence the columns c_2, c_3 , on different cycles (their Lemma 2.3(1)). In either case the final switch involves the rows r_3, x , both outside $\{1, 2\} \cup R_Q$, along a row cycle avoiding c_2 ; it changes no cell of R or of $Q(L) \setminus \{(r_3, c_3)\}$ and replaces $s_3 = s_2$ by $L[x, c_3] \neq s_2$ (or $L[x, c_2]$), so $L' \in Y_2 \setminus X_3$, as in [1]. This gives at least $\sum_L g(L) \geq f|X_3|/30$ switches from X_3 to $Y_2 \setminus X_3$. Their inverse count applies unchanged: from L' the row r'' with $L'[r'', c_3] = s_2$ equals x , and in each of the two cases L is determined, by switching back $\rho_{L'}(r_3, r'', c_3)$ and, in the second case, then the (c_2, c_3) -cycle of L'' through r'' (the same set of rows as before); so each L' has at most two preimages. Hence $|Y_2 \setminus X_3| \geq f|X_3|/60$ and $p_2 = |X_3|/|Y_2| \leq 1/(1 + f/60) \leq 60/f$. $\square$

Proof of Theorem 7.1. Claim 1 twice from $p_2 \leq 60/(n(1 - \alpha))$, as in (16) of [1]: with $\phi_j(x) = (nx + 1)/(n(1 + x - 2\alpha - \beta) - 3j - 2)$, which is non-decreasing for $j = o(n)$, $p_1 \leq \phi_2(p_2) \leq (60/(1 - \alpha) + 1)/(n(1 - 2\alpha - \beta) - 8)$ and $p_0 \leq \phi_1(p_1) \leq (np_1 + 1)/(n(1 - 2\alpha - \beta) - 5)$, so $p_0 \leq D(\alpha, \beta)/n$ with $D(\alpha, \beta) = (60/(1 - \alpha) + 2 + o(1))/(1 - 2\alpha - \beta)^2 \rightarrow 62$. $\square$

Corollary 7.2 (per-cell bound and (FC) in $S(R)$). *With R, Q as in Theorem 7.1, for a row $r \notin \{1, 2\}$ (in R_Q or not) and any (c, s) such that Q has no cell $(r, c, \cdot)$, $\mathbb{P}[L[r, c] = s \mid L \supseteq R \cup Q] \leq \Delta/n$ with $\Delta = (D + 1)/(1 - 2\alpha - \beta) + o(1)$, which is ≤ 66 for α, β small. Consequently, for every partial Latin square Q on rows other than 1, 2 with $|R_Q|, |C_Q| \leq \varepsilon n$ (ε a small absolute constant), (6) holds with $C = \Delta$: $\mathbb{P}[Q \subseteq L \mid L \supseteq R] \leq (\Delta/n)^{|Q|}$.*

Proof. The first statement is Lemma 4.1 of [1] with $\mathcal{R} = \{r\}$: its switch $\eta_L(r, r_1, c)$, $r_1 \notin E \cup \{r\}$, avoids the cells of $R \cup Q$ when $c_1 \notin C_Q$ and $r_2 \notin E$ (at most $|C_Q| + |E|$ excluded choices of r_1) and changes only rows $r, r_1, r_2 \notin \{1, 2\}$; the images L' all have $r'(L') = r_1 \notin E$ (the row with $L'[r_1, c] = s$), so in their sum (19) only $r' \notin \mathcal{R} \cup E$ occur, and for these the inverse count uses Theorem 7.1 for $Q' = Q \cup \{(r', c, s), (r, c, s')\}$ with r' as its r' : the only cell of Q' in row r' is in column c , as the theorem allows, and α, β increase by $O(1/n)$. The second statement follows by conditioning on the cells of Q one at a time. $\square$

Corollary 7.3 (the ν -tail). *For every λ, α, β and every M with $2M + 2 \leq \varepsilon n$, $\mathbb{P}_X[\nu \leq M] \leq 8\Delta^{2M+2}/n$. In particular $\mathbb{P}_X[\nu \leq C \log \log n] = O((\log n)^{2C \log \Delta}/n) = o(1/\log n)$ for every C .*

Proof. Corollary 7.2 is (6) with $C = \Delta$, in every $S(R)$; the computation after (6) in §6(e) gives the bound. $\square$

So the first term of the first-pair condition is settled, and the weak conjecture follows from the second alone:

Hypothesis 7.4 (hazard hypothesis). *For some absolute C and $M = M(n) \leq C \log \log n$, $\sup_{\lambda, \alpha, \beta} \mathbb{P}_X[\nu > M, \text{ the candidates } t = 1, \dots, M \text{ are all bad}] = o(1/\log n)$.*

Corollary 7.5. *Hypothesis 7.4 implies Hypothesis 6.12, hence Conjecture 1.1. It holds with $M = \lceil C \log \log n \rceil$ if, for $t \leq M$, the conditional probability that candidate t is bad, given $\nu > t$ and that candidates $1, \dots, t-1$ are bad, is at most $1 - c$ for an absolute $c > 0$, and $C \log(1/(1 - c)) > 1$.*

Proof. $\mathbb{P}_X[G^c] \leq \mathbb{P}_X[\nu \leq M] + \mathbb{P}_X[\nu > M, t \leq M \text{ all bad}]$, and the first term is $o(1/\log n)$ by Corollary 7.3; under the conditional bound the second is at most $(1 - c)^M = (\log n)^{-C \log(1/(1 - c))}$. $\square$

Remark 7.6 (what changed, and what did not). Hypothesis 7.4 is a statement about $O(\log \log n)$ events of constant probability at column pairs selected by the frame-selected row pair; it is a lower bound (each candidate is good with conditional probability $\geq c$) and is not a fixed-cell event, so Corollary 7.2 does not reach it. The obstacle of §9 is therefore unchanged in kind but reduced in scope: the upper-bound half — short structures at a frame-selected pair — is now available, through the η -switching of [1], whose forward degree is deterministic; only lower bounds at square-selected positions remain, and §8 shows which: at the first candidate, lengths at scale n are available (Proposition 8.3), good pairs across the cut between x_1 and y_1 are not. In the data (Remark 6.14) the conditional probabilities of Corollary 7.5 (that a candidate is bad) are ≈ 0.6 at every n and type, i.e. $c \approx 0.4$. Where the obstruction of our earlier analysis went: with the admissible rows of their Claim 4 taken to be all free rows off γ_L , the column-cycle switch through x fails to be legal exactly when the cycle through x contains one special row; the statistic g restricts to rows whose cycles contain none, and Claim 3' shows that these are $\Omega(n)$ on average, by cutting all active cycles at once with the cuts confined to arcs. Everything stays in $S(R)$; neither the mark nor the isotopy argument of Remark 6.14 is used. The constants ($D \rightarrow 62$, $\Delta \leq 66$, against their 22 and 23) are not optimised and are irrelevant for $M = O(\log \log n)$. The argument was checked by two independent automated referees (the first found a claim stated per arc that is true only in aggregate, corrected above) and numerically on sampled completions at $n = 30, 50, 100$: every restricted cross-switch, cut and Claim-4' move behaved as stated, the aggregate of Claim 2' held for every arc size with a factor about 2 to spare, and the mean of g/f was 0.22–0.25 against the bound $1/30$ [14].

8 The hazard term: an exact identity, the first candidate, and fixed positions

This section records what we can prove about Hypothesis 7.4 with the tools of §7: an exact switching identity for the conditional hazard, which locates the difficulty on both of its sides; a lower bound, at scale n , on the length of the column cycle and of its arcs at the first candidate, conditionally on x_1, y_1 sharing a cycle there (the joined part of “bad”); and the statement that

the same machinery *does* give at fixed positions, which shows exactly what the selection of the candidate columns by the square costs. None of it proves the hypothesis at any candidate, and §8(d) says why the first candidate is not reached. Throughout, R and the mark are fixed and all probabilities are in $S(R)$; the statements for $\mathbb{P}_X$ follow by averaging. The special rows are the frame rows 1, 2 and x_1, y_1 ; a row is *free* if it is not special; τ_t , $Z_t(\cdot)$ and “good” are as in §6(e), and $D_t = Z_t(x_1)$, its rows listed cyclically reading from column q_t as in §7.

(a) The switching identity

For $j \geq 1$ let $E_j = \{\nu > j, \text{ candidates } 1, \dots, j \text{ bad}\}$ and $F_j = \{\nu > j, 1, \dots, j-1 \text{ bad, } j \text{ good}\}$, and let $h_j = \mathbb{P}[F_j]/(\mathbb{P}[E_j] + \mathbb{P}[F_j])$ be the conditional hazard at j . A bad candidate is *joined* if $y_1 \in D_j$ and *separated-inadmissible* otherwise. For $L \in E_j$ with j joined, a *toggle* is a row switch $\rho_L(z, z', c_j)$ of two free rows $z, z' \in D_j$ which interleave x_1, y_1 in the cyclic order of D_j (exactly one of x_1, y_1 lies strictly between them) and whose footprint, the row cycle of (z, z') through c_j , avoids q_j . By [1, Lemma 2.4(1)] it separates x_1 from y_1 on the (q_j, c_j) -cycles; it changes only rows z, z' , so it stays in $S(R)$ and leaves ρ_{x_1, y_1} , ν and every candidate unchanged. It is *ok* if candidate j is good afterwards (separated and admissible) and *preserving* if moreover the result lies in F_j (candidates $1, \dots, j-1$ still bad). For $L' \in F_j$ a *merge* is a row switch $\rho_{L'}(z, z', c_j)$ with z a free row of $Z_j(x_1)$ and z' a free row of $Z_j(y_1)$ whose footprint avoids q_j ; it joins the two cycles with z, z' at the junctions ([1, Lemma 2.4(2)]), and it is *preserving* if the result lies in E_j . Let $f_j(L)$ and $b_j(L')$ be the numbers of preserving toggles and merges ($f_j = 0$ when j is separated-inadmissible).

Proposition 8.1 (the hazard identity). *For every $j \geq 1$ with $\mathbb{P}[E_j]\mathbb{P}[F_j] > 0$,*

$$\frac{\mathbb{P}[E_j]}{\mathbb{P}[F_j]} = \frac{\mathbb{E}[b_j \mid F_j]}{\mathbb{E}[f_j \mid E_j]}, \quad b_j \leq |Z_j(x_1)| |Z_j(y_1)| \leq n^2/4.$$

In particular $h_j \geq c$ iff $\mathbb{E}[b_j \mid F_j] \leq \frac{1-c}{c} \mathbb{E}[f_j \mid E_j]$.

Proof. A preserving toggle of L at (z, z') is undone by the row switch $\rho_{L'}(z, z', c_j)$ ([1, Fact 2.1]: the same set of cells), which is a preserving merge of L' with $z \in Z_j(x_1)$, $z' \in Z_j(y_1)$; conversely a preserving merge of L' at (z, z') is undone by a preserving toggle of its result at the same pair (the merged cycle has z, z' at the junctions, hence interleaving, and the footprint is the same row cycle, avoiding q_j ; a switch whose footprint contained q_j would conjugate the column permutation and change no candidate’s status). So the pairs $(L, \text{preserving toggle})$ and $(L', \text{preserving merge})$ are in bijection, and double counting gives $|E_j| \mathbb{E}[f_j \mid E_j] = |F_j| \mathbb{E}[b_j \mid F_j]$. $\square$

Remark 8.2 (what the identity shows). Both sides are counts of local moves inside $S(R)$, so the hazard at j is a comparison of two cycle statistics in two conditioned ensembles. With the trivial bound $b_j \leq n^2/4$ one needs $\mathbb{E}[f_j \mid E_j] \geq c'n^2$: the arcs of D_j between x_1 and y_1 must have length $\Theta(n)$, and a positive density of the pairs across them must be good and preserving. In the data (Remark 6.14; [14], `s13_hazard.py`) the preserving fraction decays like $(1 - \varepsilon)^{j-1}$ with $\varepsilon \approx 0.05$ at $n \leq 100$ (an ok toggle flips an earlier candidate s when its footprint, of length $\approx 0.35n$, meets exactly one of q_s, c_s and the pair interleaves x_1, y_1 on D_s), so that the bound the chain of conditional hazards gives on $\mathbb{P}[E_M]$ with the trivial backward bound stays above a positive constant for every M (if $\varepsilon(n)$ stays bounded below, which the data do not decide); the identity forces $\mathbb{E}[b_j \mid F_j]$ to decay at the same rate, and it does (the measured ratio of the two sides is ≈ 1.5 at $j \leq 6$, matching $\mathbb{P}[E_j]/\mathbb{P}[F_j]$ for $j \leq 4$). So what fails is not the move

family but the two obvious bounds on its two sides; a direct comparison of the two conditional expectations is not excluded, and we do not have one. The data are for the type (n) only.

(b) Lengths at the first candidate

Let $E_1^j = \{x_1 \sim y_1 \text{ on } D_1\}$ (the joined part of E_1 ; $\nu \geq 2$ always by Lemma 6.9), $\ell = |D_1|$, and a, a' with $a + a' = \ell$ the numbers of rows of the two arcs of D_1 between x_1 and y_1 (each arc counted with one of its endpoints). A (q_1, c_1) -cycle is *inactive* if it contains no special row, and $g(L)$ is the number of rows on inactive cycles.

Proposition 8.3 (cycle and arc lengths at the first candidate). *Let $n \geq \Delta^{32}$, $8 \leq k \leq n/100$ and $1/n \leq \delta \leq 1/100$. Then*

$$\mathbb{P}[\ell \in [k, 2k] \mid E_1^j] \leq 400 \frac{k}{n}, \quad \mathbb{P}[\min(a, a') \in [k, 2k], \ell \geq \delta n \mid E_1^j] \leq 6000 \frac{k}{\delta^2 n},$$

and consequently, with $k_0 = \lfloor \log n / (4 \log \Delta) \rfloor$,

$$\mathbb{P}[\ell \leq \delta n \text{ or } \min(a, a') \leq \delta^3 n \mid E_1^j] \leq 48000 \delta + \frac{5(\Delta^5 + k_0 \Delta^2)}{\sqrt{n} \mathbb{P}[E_1^j]}.$$

The mechanism is that of Claim 4': $\Omega(kn)$ moves out of the class, $O(k^2)$ into any square, the inverse pinned by the two marked rows x_1, y_1 , which must both lie in a short piece; the dyadic sum $\sum_k k/n$ over $k \leq \delta n$ is $O(\delta)$, which is the right order ($\mathbb{P}[\ell \leq \delta n] \approx \delta$ for a random permutation). Nothing here bounds a specific cycle from below; the proposition says that a cycle through two marked rows is unlikely to be of intermediate length. The constants are not optimised and the statement is purely asymptotic ($n \geq \Delta^{32}$ comes from matching the dyadic range with the fixed-cell range at $k_0 \geq 8$). If $\mathbb{P}[E_1^j]$ is small the first-candidate hazard is large anyway, so the denominator is harmless.

Proof. (i) The cycle length. Fix k and let $\mathcal{C}_k = E_1^j \cap \{\ell \in [k, 2k]\}$. *Moves.* For $L \in \mathcal{C}_k$, a free row $z \in D_1$ and a row z' on an inactive cycle C : if $\rho_L(z, z', c_1)$ avoids q_1 , switch it; otherwise first switch the column cycle C (columns q_1, c_1 , rows of C) to obtain L'' , in which $\rho_{L''}(z, z', c_1)$ avoids q_1 ([1, Lemma 2.3(1)] after re-basing the row cycle at a third of its columns, which exist since z, z' lie on different (q_1, c_1) -cycles; directly as in the proof of Claim 4', since $C \not\ni z$), and then switch it. The column switch changes only rows of C , which are not special; the row switch only rows z, z' . So rows $1, 2, x_1, y_1$ are unchanged and the result L' is in $S(R)$ with the same mark, first pair, ρ_{x_1, y_1} , ν and candidates. By [1, Lemma 2.4(2)] D_1 and C merge into one cycle $D' = D_1 \cup C$ and every other cycle is unchanged; directly, column q_1 is unchanged and the c_1 -entries of z, z' are exchanged, so the column permutation is composed with the transposition of $L[z, c_1], L[z', c_1]$, which splices C into D_1 immediately after z in the list. Hence $L' \in E_1^j$ with $\ell' = \ell + |C| \geq \ell + 2$ and the cyclic order of the special rows on the cycle preserved. There are at least $(\ell - 4)g(L)$ moves from L . *Inverse count.* If L' arises from (L, z, z') then L is recovered from (L', z, z') by switching $\rho_{L'}(z, z', c_1)$ ([1, Fact 2.1]) and then, in the two-step case, switching the (q_1, c_1) -cycle through z' ; so L' has at most $2N$ preimages, N the number of unordered pairs $\{z, z'\}$ of free rows on D' such that switching $\rho_{L'}(z, z', c_1)$ splits D' into a piece $P \ni x_1, y_1$ with $|P| \in [k, 2k]$ and a piece S . By [1, Lemma 2.4(1)] S consists of the rows strictly between z and z' in the list together with z' , so $P = D' \setminus S$ is a contiguous arc of D' ending at z , which determines $\{z, z'\}$; and a contiguous arc of length m containing x_1 can be

placed in at most m ways. Hence $N \leq \sum_{m=k}^{2k} m \leq 3k^2$ and L' has at most $6k^2$ preimages. *Lower bound on g .* The proof of Claim 3' applies verbatim with special set $\{1, 2, x_1, y_1\}$, $Q = \emptyset$ and the class $\mathcal{C}_k$: its moves (row switches of good same-arc pairs on the cycles through rows 1, 2 other than D_1 , and the cross-switches of Claim 2' on those cycles) touch neither D_1 nor a special row, so ℓ , E_1^j , the mark and the candidates are preserved. The free rows off D_1 number at least $n - 2k - 4$ and lie on the switched active cycles or on inactive cycles; the forward count is at least $\frac{1}{27}((n - 2k - g - 13)_+)^2$ in aggregate, the backward count at most $ng(L')$, and Jensen's inequality gives $\mathbb{E}[g \mid \mathcal{C}_k] \geq (n - 2k)/31$ for $k \leq n/100$ and n large. *Conclusion.* $|\mathcal{C}_k| (k - 4) \mathbb{E}[g \mid \mathcal{C}_k] \leq 6k^2 |E_1^j|$, so $\mathbb{P}[\ell \in [k, 2k] \mid E_1^j] \leq 6k^2 \cdot 31 / ((k - 4)(n - 2k)) \leq 400k/n$ for $8 \leq k \leq n/100$.

(ii) *The arcs.* Let $\mathcal{C}'_k = E_1^j \cap \{a \in [k, 2k], \ell \geq \delta n\}$, a the arc from x_1 to y_1 in the list order (the other arc by symmetry). If $k > \delta n/4$ or $\delta n < 240$ the bound is ≥ 1 ; assume otherwise. The other arc has $\ell - a \geq \delta n/2$ rows and at most two frame rows, so it contains a special-row-free stretch of at least $\delta n/6 - 1$ free rows; let $B = B(L)$ be the longest (ties by position), $b = |B|$. *Transfer moves.* For a good same-arc pair (w, w') in B (footprint avoiding q_1) and a free row z of the x_1 - y_1 arc (at least $k - 3$ choices): switch $\rho_L(w, w', c_1)$, which by [1, Lemma 2.4(1)] cuts the rows strictly between w and w' together with w' off as a new cycle S containing no special row, hence inactive, leaving x_1, y_1 and their arc untouched; then perform the insert move of (i) at the pair (z, z') with $z' := w'$. The result is in E_1^j with $\ell' = \ell$, $a' = a + |S| \geq a + 2$, the same cyclic order of special rows, and S spliced in after z and ending at w' (in the two-step case the column switch reverses S as a cycle and the splice still ends at $z' = w'$). Claim 2', applied to the class of triples $(L, D_1, B(L))$ with $L \in \mathcal{C}'_k$ and $|B(L)| = b$ — closed under its cross-switches, which act inside B , touch no special row, keep B as a set ([1, Remark 2.7]) and change neither ℓ nor a — gives at least $(b - 3)^2/9 \geq \delta^2 n^2/400$ good pairs in B in aggregate, hence at least $(k - 3)\delta^2 n^2/400$ moves per square of $\mathcal{C}'_k$ on average. *Inverse count.* From L' : the segment S to be cut out of the x_1 - y_1 arc leaves a prefix and a suffix of that arc of total length in $[k, 2k]$, at most $(2k)^2$ choices, and determines (z, z') (z the row before S , $z' = w'$ its last row); undoing the insert is the row switch at (z, z') and, in the two-step case, the column switch of the cycle through z' (two alternatives); undoing the cut is the row switch at (w, w') with w' known and w one of at most n rows. So at most $8k^2 n$ preimages, and $\mathbb{P}[a \in [k, 2k], \ell \geq \delta n \mid E_1^j] \leq 8k^2 n \cdot 400 / ((k - 3)\delta^2 n^2) \leq 6000k/(\delta^2 n)$ for $k \geq 8$.

(iii) *Short lengths.* $\{\ell = m\} \cap E_1^j$ is the union, over the positions of x_1, y_1 (n^2), of q_1, c_1 (n^2), of the $m - 2$ other rows of the cycle (n^{m-2}), of the position of y_1 on it ($m - 1$) and of the $m - 1$ free symbols (n^{m-1} ; $L[y_1, q_1] = L[1, p']$ is forced), of fixed-cell events with $2m + 2$ cells (the cycle and the anchors $(x_1, p), (y_1, p)$); by Corollary 7.2 and the exchangeability of rows $3, \dots, n$, as in Corollary 7.3, its probability is at most $(m - 1)\Delta^{2m+2}/n$. For $\{a = m\} \cap E_1^j$ the path from x_1 to y_1 has $m + 1$ rows and $2m + 2$ cells, and c_1 is pinned by the further cells (x_1, p') and (y_1, c_1) : $2m + 5$ cells, n^{2m+4} placements, probability at most Δ^{2m+5}/n . Summing over $m \leq k_0$ (at most $2k_0\Delta^2/\sqrt{n}$ and $4\Delta^5/\sqrt{n}$, using $\Delta^{2k_0} \leq \sqrt{n}$), (i) over the dyadic $k \in [k_0, \delta n]$ ($\leq 800\delta$) and (ii) over the dyadic $k \in [k_0, \delta^3 n]$ for each arc ($\leq 2 \cdot 12000\delta$) gives the consequence. $\square$

Remark 8.4. The same argument, with the frame rows (when they lie on D_1) as the pinning rows, bounds the lengths of the stretches of D_1 between consecutive special rows below by $\delta^3 n$ with probability $1 - O(\delta)$; and it applies to D_j for any $j < \nu$ as a statement about $\mathbb{P}[\cdot \mid x_1 \sim y_1 \text{ at } j]$, but not about $\mathbb{P}[\cdot \mid E_j]$, since its row switches have footprints of length $\Theta(n)$ that change the status of earlier candidates. On sampled completions at $n = 30, 50, 100$ every sampled insert and transfer move behaved as stated and was undone by the stated inverse, the

preimage counts N were at most $3k^2$, $\mathbb{E}[g/(n-\ell)]$ was at least 0.26 in every bin of ℓ/n (bound $1/31$), and $\mathbb{P}[\ell \in [k, 2k) \mid E_1^j]$ tracked the permutation value $3k^2/n^2$ [14].

(c) Fixed positions

Proposition 8.5 (two fixed rows at a fixed column pair). *Let R be a $2 \times n$ Latin rectangle (rows $1, 2$), $n \geq 504$, $q \neq c$ columns and $x \neq y$ rows in $\{3, \dots, n\}$. Then*

$$\mathbb{P}_{S(R)}[x \text{ and } y \text{ lie on the same } (q, c)\text{-column cycle}] \leq 1 - \frac{1}{30} + \frac{3}{n}.$$

For a uniform Latin square of order $n \geq 502$ and arbitrary rows $x \neq y$ the bound is $1 - \frac{1}{30}$.

Proof. Let D_x be the (q, c) -cycle through x and $\ell_x = |D_x|$. For a permutation τ of $\{3, \dots, n\}$ fixing x , $L \mapsto \tau L$ is a bijection of $S(R)$ with $y \in D_x(L)$ iff $\tau y \in D_x(\tau L)$; averaging over τ , $\mathbb{P}[y \in D_x] = \mathbb{E}[|D_x \cap \{3, \dots, n\} \setminus \{x\}|/(n-3)] \leq \mathbb{E}[\ell_x - 1]/(n-3)$. The proof of Claim 3' applies verbatim with special set $\{1, 2, x\}$, $Q = \emptyset$ and the class $S(R)$ (at most three arcs; none of its moves touches a special row), and gives $\mathbb{E}[g] \geq f/30$ for the number g of rows on cycles containing no special row, $f = n - 3 \geq 501$. Since $\ell_x \leq n - g$, $\mathbb{E}[\ell_x - 1] \leq n - 1 - (n-3)/30$. For a uniform square take the special set $\{x\}$, $f = n - 1$, and average over all τ fixing x . $\square$

Corollary 8.6. *For a uniform Latin square and a fixed column pair, with M the length of the longest cycle of the pair, $\mathbb{E}[n - M] \geq (n-1)/60$ and $\mathbb{P}[n - M \geq n/120] \geq 1/120 - O(1/n)$; by transposition the same holds for the longest cycle of a fixed row pair.*

Proof. $\sum_C |C|(n - |C|) = \sum_x (n - \ell_x) \leq 2n(n - M)$, and $\mathbb{E} \sum_x (n - \ell_x) \geq n(n-1)/30$ by the proposition. $\square$

This is a statement in the “coarse” direction that the surviving material of [3] does not give ([6], §4): with probability bounded below, the rows off the longest cycle of a column pair are a positive fraction of all rows. It is not an improvement for the type (n) itself, for which the intercalate argument of [6] gives $5/6 + o(1)$, and it does not approach $\mathbb{P}[(n)] = o(1)$.

(d) What the first candidate still needs

Proposition 8.5 is the first-candidate statement with (q, c) and y fixed: it needs no pair across any cut, because y is exchangeable with the other free rows and “is y on D_x ” becomes “how long is D_x ”, which Claim 3' answers. In the ladder, y_1 is selected by the cell (y_1, p) and the candidate columns q_1, c_1 — the columns of row y_1 holding $L(1, p')$ and $L(x_1, p')$ — are functions of row y_1 ; conditioning on the columns destroys the exchangeability of y_1 , and conditioning on less leaves the columns random. So by Proposition 8.1 with $j = 1$ and $b_1 \leq n^2/4$, the first-candidate hazard $h_1 \geq c$ follows from $\mathbb{E}[f_1 \mid E_1] \geq c'n^2$, and Proposition 8.3 with the remark after it supplies the lengths ($\geq \delta^6 n^2$ interleaving pairs whose split is admissible, with probability $1 - O(\delta)$ on E_1^j) but not the goodness: a positive density of the pairs across the x_1 - y_1 cut must have footprint avoiding q_1 . This does not follow from Claim 2': the cross-switch at a bad interleaving pair reverses the segment between the two rows containing one of x_1, y_1 and no frame row (legal, and it preserves ν and every candidate as an unordered pair), and by [1, Lemma 2.6] it toggles the goodness of every pair with exactly one row in that segment — but it toggles their interleaving status as well, since the inner row and the inner special row exchange their order; so the Kind-2 image of a bad interleaving pair is a good *non-interleaving* pair, outside the family, and with

same-stretch auxiliary pairs the Kind-1 image leaves the family instead. More than that — and this whole paragraph is heuristic, not a theorem — the configuration in which a pair of free rows of D_1 at distance ≥ 2 is good iff it does not interleave x_1, y_1 (adjacent pairs $\{z, \tau_1(z)\}$ are always bad), with inactive cycles carrying an orientation-dependent affinity to one arc, is preserved, up to the $O(n)$ pairs containing a switched row, by every move used in this section and in §7 — cross-switches, inserts, cuts, transfers and column switches of inactive cycles — checked move by move, not proved; and on it $f_1 = 0$. We therefore see no counting argument from these moves that gives the goodness across the cut. (The truth is the opposite: in the data the good fraction is 0.50 among interleaving and among non-adjacent same-arc pairs alike, at every distance ≥ 2 along D_1 , with adjacent pairs never good [14].) The separated-inadmissible part of E_1 (15–17% of the bad candidates in the data, pooled over j), on which $f_1 = 0$, needs a toggle that moves a frame row off the cycle of x_1 , which we have not analysed. So the obstacle of §9 has, at the first candidate, the exact form: a positive density of good pairs across the x_1 – y_1 cut, without the exchangeability of y_1 ; and at later candidates, in addition, the preservation of the earlier ones.

9 What remains

Hypotheses 5.1, 6.4, 6.7, 6.12 and 7.4 are statements about configurations at positions selected by the frame: a short ρ_{x_k, y_k} -cycle through p avoiding p' ; a clean column pair separated by the mark; a short clean column cycle at a prescribed column pair; a column pair along ρ_{x_1, y_1} whose column permutation separates x_1 from y_1 ; and, in the last, the same at one of the first $O(\log \log n)$ such pairs, given that the earlier ones fail. Their heuristic frequencies are, for large n , far above what is needed. The weakest of the sufficient conditions our identities provide (short of EQ(δ) itself) is the disjunction of (L) and Hypothesis 7.4 class by class,

$$\sup_{\lambda, \alpha, \beta} \min(\mathbb{P}_X[\text{Flip} = \emptyset], \mathbb{P}_X[\nu > M, t \leq M \text{ all bad}]) = o(1/\log n) \quad \text{for some } M = M(n) \leq C \log \log n$$

(by Theorems 2.4, 4.2 and 3.3, Proposition 6.11 and Corollary 7.3; by Remark 3.4 the supremum may be restricted to types with at most $\lceil 16 \log n \rceil$ parts), which we record but do not promote, since it is only a disjunction. Two features make all of these statements hard for the available tools: the positions are selected by the frame of the mark, and the whole space is conditioned on the type of rows 1, 2. The first is removable (by row exchangeability a frame-selected row pair is two prescribed cells, see below); the second is not, but §7 shows that for *upper* bounds on events at a sparse set of cells (occupying at most εn rows and columns) it can be handled by the switchings of [1], and §8 shows that the same switchings give *lower* bounds at square-selected positions for cycle *lengths* (Proposition 8.3) and give the different-cycles bound itself at *fixed* positions, through row exchangeability (Proposition 8.5); what no tool of ours reaches is a lower bound for a different-cycles event at square-selected positions, where the second marked row is not exchangeable. The minimal instance of the obstacle is a lower bound on a single frame-selected pair:

$$\mathbb{P}_X[(x_1, y_1) \text{ flippable}] \geq c, \quad x_1 = \pi^{-1}(1), y_1 = \pi^{-1}(2), \quad (7)$$

which no identity of this paper gives unconditionally: flippability of a ladder pair is a class invariant of the ladder trades, and Proposition 6.10 reduces (7) one way to $\mathbb{P}_X[G] \geq 2c$, a lower bound for an event of the same kind with rows and columns exchanged (the column permutation of a pair of columns selected from rows x_1, y_1 separating the frame-selected rows

x_1, y_1), which is not known to be easier. By row exchangeability (7) is a statement at fixed positions under conditioning: it is the probability that p, p' lie in different cycles of $\rho_{3,4}$ given rows 1, 2 and the two cells $L(3, p) = L(1, p')$, $L(4, p) = L(2, p')$. Even without the conditioning on rows 1, 2 we know of no lower bound in the literature for the related quantity $\mathbb{P}[\text{the } \rho_{3,4}\text{-cycle through } p \text{ has length } k \text{ and avoids } p'], 3 \leq k \leq \log^5 n$; the one case we know to be settled is $k = 2$, where the concentration of the number of intercalates at $n^2/4$ [11] and the exchangeability of row pairs and of columns give $\approx 1/n$. We record what the available tools do and do not give.

Unconditional facts. For two generic rows of a uniform square, Proposition 3.2(i) gives, by Markov's inequality on the number of columns on short cycles, $\mathbb{P}[\text{all cycles of } \rho_{x,y} \text{ have length } \leq \varepsilon n] \leq 2\varepsilon/(1 - \varepsilon)$; so a generic pair has a cycle longer than $n/5$ with probability $\geq \frac{1}{2}$, and the column-cycle analogue holds by transposition. Kwan and Sudakov's bound $\mathbb{P}[N_2 \geq t] = e^{-\Omega(t \log t)}$ on the number of 2-cycles of two fixed rows [10, Thm. 3] is proved by joining switchings of the same kind as those of Proposition 3.2 and is likewise an upper bound on short cycles. The number of completions of any $2 \times n$ rectangle is within $e^{O(n \log^2 n)}$ of any other (van der Waerden and Brégman; [13, Prop. 4], as used in [12]), so any property of uniform squares with failure probability $e^{-\omega(n \log^2 n)}$ — such as the intercalate concentration of Kwan, Sah and Sawhney [11] — holds conditionally on any fixed rows 1, 2 and mark, hence in X for every λ, α, β ; but the natural failure probability of a statement about $m \leq n$ prescribed row pairs is $e^{-\Omega(m)}$, which this transfer does not reach.

Why the fixed-matching version is harder. With a fixed matching the orbit weight $\bar{y}_k$ of Proposition 6.3 needs a lower bound on the probability that a generic column pair is interleaved with $\{p, p'\}$ in ρ_{x_k, y_k} — an instance of the statement “two given pairs of lines of a random Latin square are interleaved in the cycle structure of a third pair with probability in $[c, 1 - c]$ ” inside X . Without the conditioning its upper half is elementary and its lower half reduces to macroscopic cycles (by conjugation invariance the interleaving probability given the cycle type of $\rho_{x,y}$ is $\frac{1}{3} \sum_i (m_i)_4 / (n)_4 + 2 \sum_{i \neq j} (m_i)_2 (m_j)_2 / (n)_4 \leq \frac{1}{3} + O(1/n)$, and $\geq \frac{1}{3} \mathbb{P}[W \mid \text{type}]^3 - O(1/n)$ by Jensen, W the event that two given columns share a cycle). Inside X the exact tools are the free trades of rows other than 1, 2 (which preserve interleaving) and the legal column turns (which are legal exactly off the interleaved configurations), and every identity they produce is fair: it equates two expectations with random weights whose lower bounds are again of the same form. Lower bounds on “different cycles” events are, in the language of [3], of the type of the direction of their Lemma 3.12 whose proof has the gap: at square-selected positions they require a repair of the same-cycle configurations with a rigid inverse, and the repairs available in X (trades or cross-switches of a free pair of third rows) have unbounded multiplicity; at fixed positions the multiplicity is absorbed by the exchangeability of the second row (Proposition 8.5), which is exactly what the selection by the square removes (§8(d)). The one case in which we found a rigid inverse is CGW's own case 3 with $\min(\alpha, \beta) = 2$, which gives $\mathbb{P}_X[A] \geq \frac{1}{4}$ for marks at distance 2 (Remark 3.5); it does not extend to other distances for a reason of type rather than multiplicity. The adaptive coordinates sidestep the question: the coordinate is placed where it is separated by construction, and the price is paid in rarity rather than in a constant we cannot prove; at the first ladder pair alone (§6(e)) the rarity disappears as well, and what is left is an upper bound on a short arc at the frame-selected pair, which §7 proves, together with a constant conditional hazard over $O(\log \log n)$ candidates along ρ_{x_1, y_1} , Hypothesis 7.4 — the simplest sufficient condition we have (not proved weaker than the others), and still a lower bound at square-selected positions.

The fixed-position form, and the literature. Row exchangeability turns every statement at the frame-selected pair into one at fixed positions: for an event E defined through the pair, $\mathbb{P}_X[E] = (n-2)(n-3)\mathbb{P}_X[E']$, where E' is E read at rows 3, 4 together with the two anchor cells $L(3, p) = L(1, p')$, $L(4, p) = L(2, p')$ that make $(3, 4)$ the selected pair. The ν -tail of the first-pair hypothesis (through (6)) is then a constant-factor upper bound on the probability of events on $O(\log \log n)$ cells in two fixed rows of a uniform completion of a fixed $2 \times n$ rectangle with two further cells prescribed — in the language of recent work on fixed-cell probabilities, a “spread” bound under dense conditioning; this is what §7 proves. The hazard hypothesis and the witness lemma involve column cycles through all rows and are not fixed-cell events. The published results stop short of dense conditioning. For uniform $k \times n$ Latin *rectangles* the one-cell bound $\mathbb{P}[L(r, c) = s \mid Q \subseteq L] \leq 1/n + O((k + |Q_r|)/n^2)$ (their proof gives $1/(n - 2(k-1) - |Q_r|)$) holds for any completable Q with $2k + |Q_r| \leq n/2$ [12, Lemma 3.3] — which allows Q to contain full rows, so (6) holds, with $C = 1 + O(k/n)$, for the first $k \leq n/4 - M - 1$ rows of a uniform rectangle extending R ; but the first k rows of a uniform completion of R to a square are not a uniform rectangle (they are weighted by their number of completions), and the two laws are known to agree only within $e^{O(n \log^2 n)}$. For $k < (\frac{1}{2} - \varepsilon)n$ rows, Divoux, Kelly, Kennedy and Sidhu [5, Thm. 1.6] prove, for every $\varepsilon, \eta > 0$ and some $\gamma = \gamma(\varepsilon, \eta) > 0$, the two-sided bound $((1 \pm \eta)/n)^{|P|}$ for γn -sparse P (at most γn cells in each row, column and symbol), and note that the sparsity is essentially necessary; in full squares the corresponding two-sided bound $(\delta/n)^{|P|} \leq \mathbb{P}[L \supseteq P] \leq (\Delta/n)^{|P|}$, $\delta \rightarrow 1/23$, $\Delta \rightarrow 23$, is proved by Allsop and Morris [1, Thm. 1.1] for P occupying at most αn rows and βn columns with $2\alpha + \beta < 1$ (or any two of rows, columns, symbols). Our $R \cup Q$ occupies all n columns and all n symbols, so none of these statements applies as stated; §7 shows that the Allsop–Morris *proof* does go through with two complete rows, once its column-cycle and cross-switches are kept away from those rows. What makes their η -switching the right tool is that it displaces the target cell within its row and otherwise changes only two other rows: a plain trade of row r with a free row along a row cycle is rigid with the same deterministic forward degree, but it rewrites row r at every column of the cycle, so it gives one factor $1/n$ per row and nothing for a second cell in the same row; the η -switch changes row r only at c and c' , which is what makes a second factor possible, at the price of the multiplicity n in the intercalate case, which their Theorem 3.1 controls (the turns of §6, by contrast, have random forward degrees, which is where the lower tails come from). The difficulty of counting column cycles that avoid prescribed rows is noted, in a different context, by Allsop and Wanless [2, §3.1]; in §7 it is sidestepped rather than solved, by cutting the cycles so that the prescribed rows are never moved. We mention finally that the parity marginal of Conjecture 1.1 follows from Kwan, Petrova and Sawhney’s work: their law of large numbers for the number of odd rows [9, Thm. 1.3(1)] and row exchangeability give $\mathbb{P}[\text{rows } 1, 2 \text{ have the same parity}] = \frac{1}{2} + o(1)$, i.e. $P_n(\sigma \text{ even}) = \frac{1}{2} + o(1)$, as for uniform derangements. (Their written deduction of Theorem 1.3(1) passes through their total-variation statement 1.3(4) and so through [4], which depends on the gap of Remark 3.5; but their large deviation principle 1.3(5) is deduced from their Theorem 6.4 and Lemma 6.5(ii) without [4], and it implies the law of large numbers, so the parity marginal does not depend on the gap.) This says nothing about any individual adjacent pair λ, μ (which have opposite parities). Their method — a canonical, switching-invariant family of “stable intercalates” that meets every set of βn rows, columns or symbols simultaneously, transferred to Latin squares by an approximation lemma for the triangle-removal process — controls events at all *large* sets of rows at once; it does not by itself say anything about a single square-selected pair, but it is the one technique

we know that reaches every large row set simultaneously, which is what a long ladder needs; we have not attempted to adapt it to completions of a fixed rectangle.

References

- [1] J. Allsop and P. Morris, Universal probability bounds for partial Latin squares, arXiv:2606.18174 (2026).
- [2] J. Allsop and I. M. Wanless, Subsquares in random Latin rectangles, *Combinatorica* 45 (2025), art. 29.
- [3] N. J. Cavenagh, C. Greenhill and I. M. Wanless, The cycle structure of two rows in a random Latin square, *Random Structures Algorithms* 33 (2008), 286–309.
- [4] N. J. Cavenagh and I. M. Wanless, There are asymptotically the same number of Latin squares of each parity, *Bull. Aust. Math. Soc.* 94 (2016), 187–194.
- [5] A. Divoux, T. Kelly, C. Kennedy and J. Sidhu, Subsquares in random Latin squares and rectangles, *J. Combin. Des.* 34 (2026), 184–197.
- [6] [author], A gap in the proof of Lemma 3.12 of Cavenagh–Greenhill–Wanless, note, October 2026.
- [7] M. J. Gill, A. Mammoliti and I. M. Wanless, Canonical labeling of Latin squares in average-case polynomial time, *Random Structures Algorithms* 66 (2025), e70015.
- [8] M. T. Jacobson and P. Matthews, Generating uniformly distributed random Latin squares, *J. Combin. Des.* 4 (1996), 405–437.
- [9] M. Kwan, K. Petrova and M. Sawhney, Parities in random Latin squares, arXiv:2509.13125 (2025).
- [10] M. Kwan and B. Sudakov, Intercalates and discrepancy in random Latin squares, *Random Structures Algorithms* 52 (2018), 181–196.
- [11] M. Kwan, A. Sah and M. Sawhney, Large deviations in random Latin squares, *Bull. London Math. Soc.* 54 (2022), 1420–1438.
- [12] M. Kwan, A. Sah, M. Sawhney and M. Simkin, Substructures in Latin squares, *Israel J. Math.* 256 (2023), 363–416.
- [13] B. D. McKay and I. M. Wanless, Most Latin squares have many subsquares, *J. Combin. Theory Ser. A* 86 (1999), 322–347.
- [14] [author], `cgw-conjecture`: scripts, logs and notes, <https://github.com/mkkinyon/cgw-conjecture>, 2026.